

# Building Image Processing Applications Using scikit-image

---

WORKING WITH IMAGE DATA



**Janani Ravi**

CO-FOUNDER, LOONYCORN

[www.loonycorn.com](http://www.loonycorn.com)

# Overview

**Working with images is increasingly important**

**Deep learning algorithms and AI rely on images and videos as input**


**scikit-image is an image processing toolkit**

**Powerful, versatile, well-implemented**

# Prerequisites and Course Outline

---

# Prerequisite Courses



**Python: Getting Started**

**Python Fundamentals**

**Working with Multidimensional Data  
Using NumPy**

# Software and Skills



**Basic understanding of Python programming**

**Python 3, NumPy, Matplotlib**

**Working with Jupyter notebooks**

**Basic knowledge of statistics**



# Course Outline

## **Working with images**

- Manipulating images as arrays of pixels
- Detecting contours
- Finding convex hull

## **Detecting features and objects**

- Extracting features
- Detecting corners
- Denoising images

## **Segmentation and transformation**

- Local and global thresholding
- Image segmentation
- Image transformations

# Introducing scikit-image

---

# scikit-image

Image processing toolbox for SciPy; high-quality, peer-reviewed code, available free of charge and free of restriction





# scikit-image

**Simple Python APIs for complex image processing tasks**

**Gallery of examples of image manipulation techniques**

**Open source community of developers**

**Part of SciPy, open source Python software for engineers**

# Image Pre-processing Methods

Uniform Aspect  
Ratio

Uniform Image Size

Mean and Perturbed  
Images


Normalized Image  
Inputs

Dimensionality  
Reduction


Data Augmentation

**Common techniques to improve performance of  
machine learning models**

# Pixels in Images



Copyright: Wikimedia commons,  
GFDL and cc-by-sa-2.5,2.0,1.0



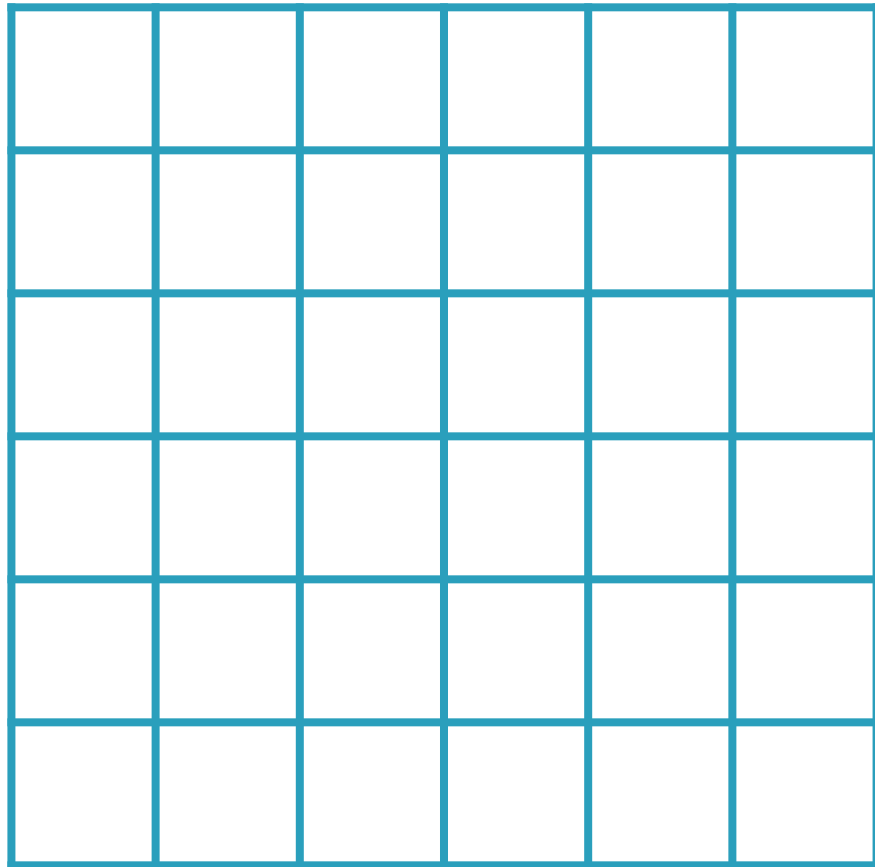
Copyright: Wikimedia commons,  
GFDL and cc-by-sa-2.5,2.0,1.0

# Images as Matrices





# RGB Images

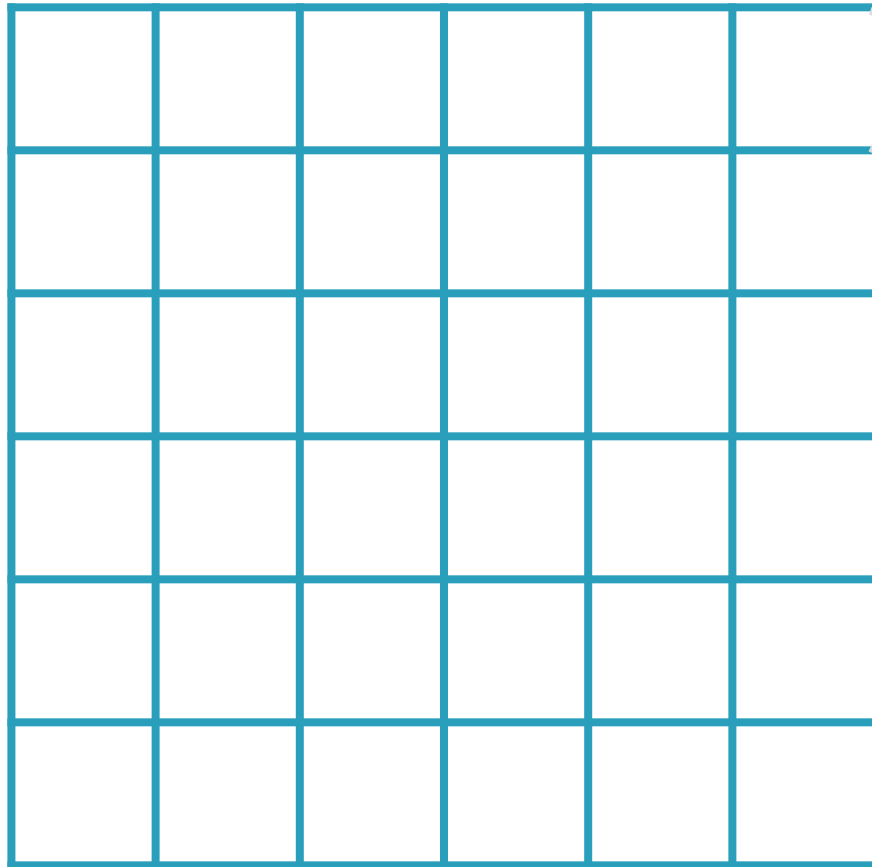


**RGB values are  
for color images**

**R, G, B: 0-255**



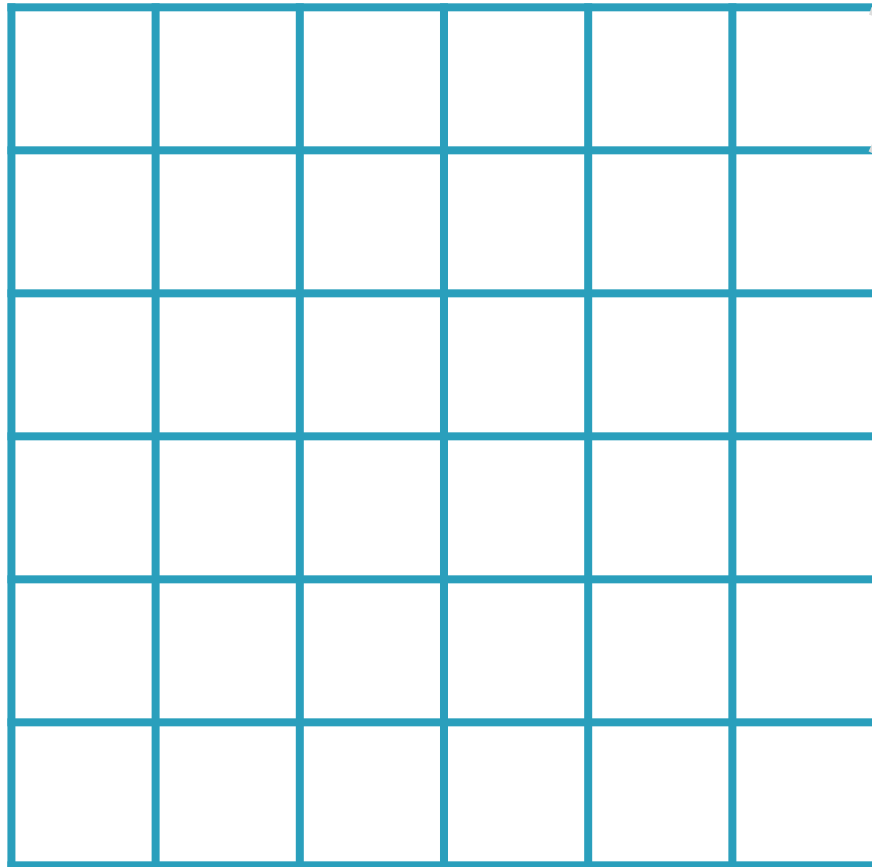
# RGB Images



**255, 0, 0**



# RGB Images

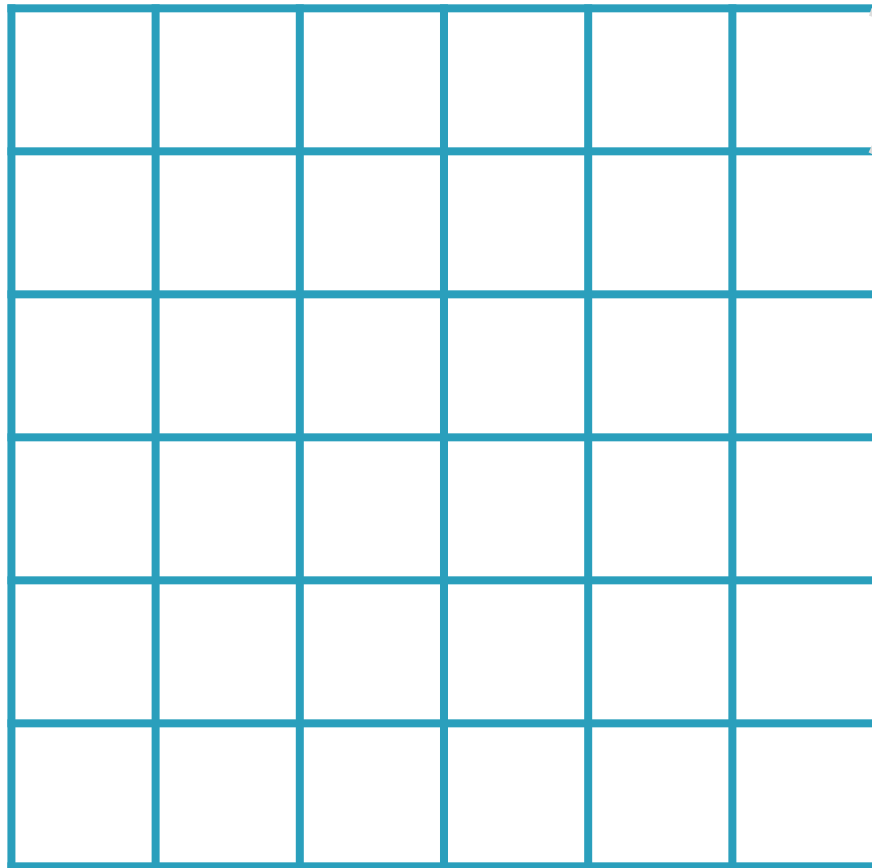


0, 255, 0





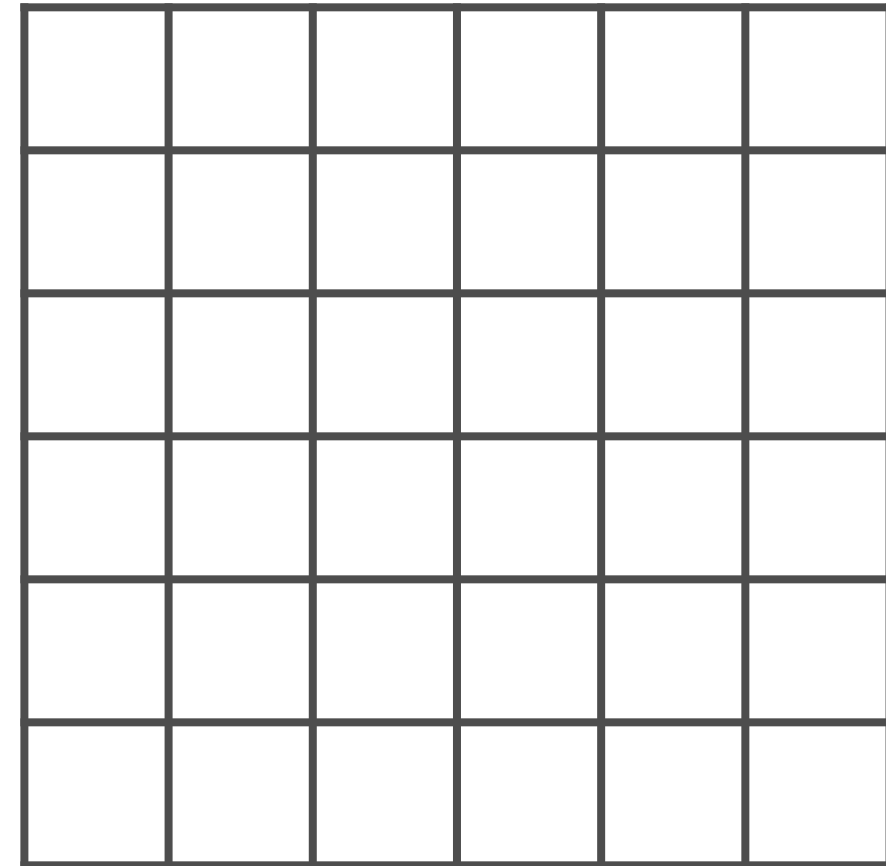
# RGB Images



0, 0, 255

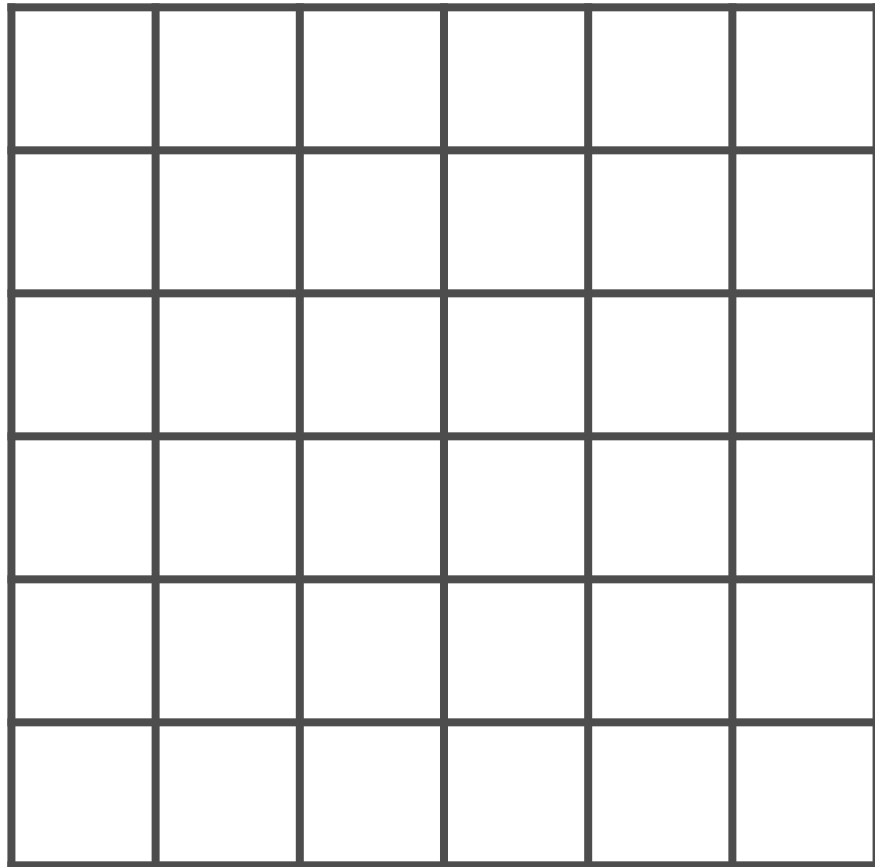
**3** values to represent  
color, **3** channels

# Grayscale Images





# Grayscale Images




**Each pixel represents  
only intensity information**

**0.0 - 1.0**



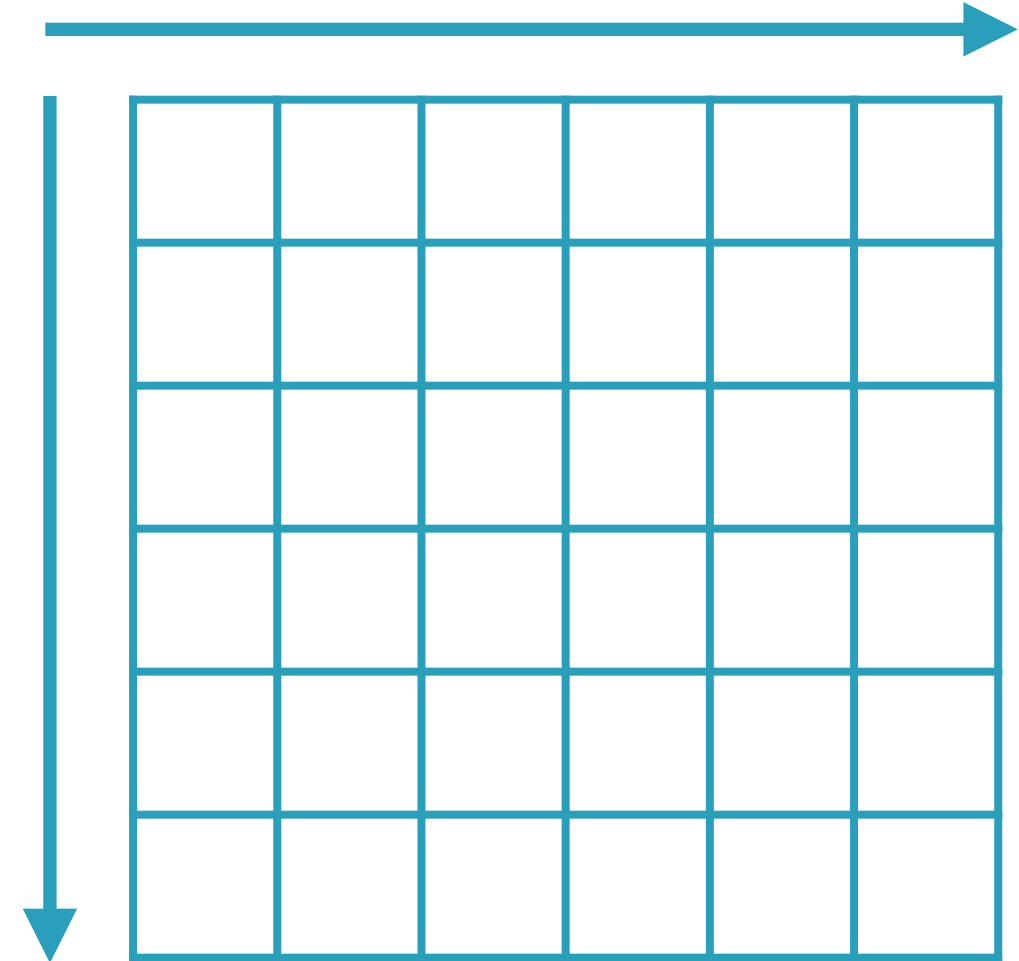
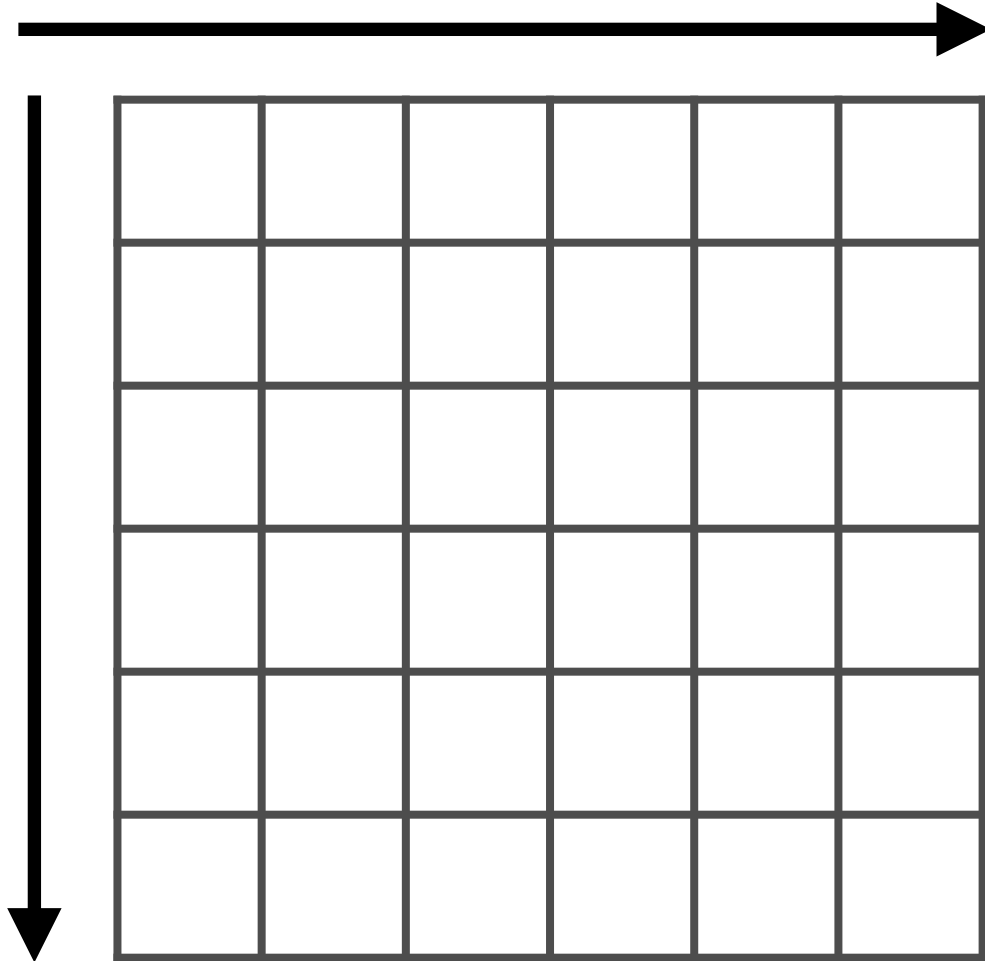
# Grayscale Images



0.5


**1** value to represent  
intensity, **1** channel

# Images as Matrices




**Images can be represented by a 3-D matrix**

# Images as Tensors



**(6, 6, 1)**



**(6, 6, 3)**

# List of Images



ML frameworks (e.g. TensorFlow) usually deal with a **list of images in one 4-D Tensor**

# List of Images



**The images should all be the same size**





List of Images

(10, 6, 6, 3)

**The number of channels**



List of Images

(10, 6, 6, 3)

**The height and width of  
each image in the list**



List of Images

(10, 6, 6, 3)

**The number of images**

Demo

**Working with images using NumPy**

# BlockView and Pooling

---

# Block Views

Extract non-overlapping image patches in order to perform local operations on these patches

# Pooling

A window function which performs an aggregation on the input to extract significant features

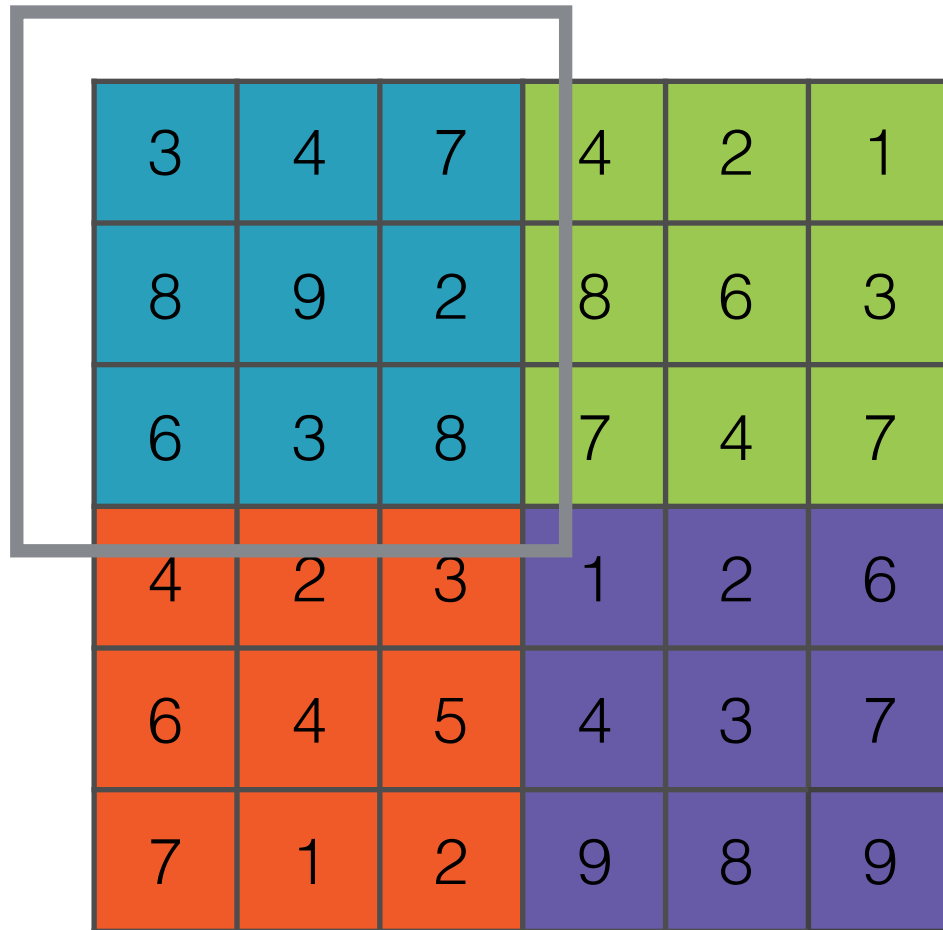
# Max Pooling

3	4	7	4	2	1
8	9	2	8	6	3
6	3	8	7	4	7
4	2	3	1	2	6
6	4	5	4	3	7
7	1	2	9	8	9

9	8
7	9

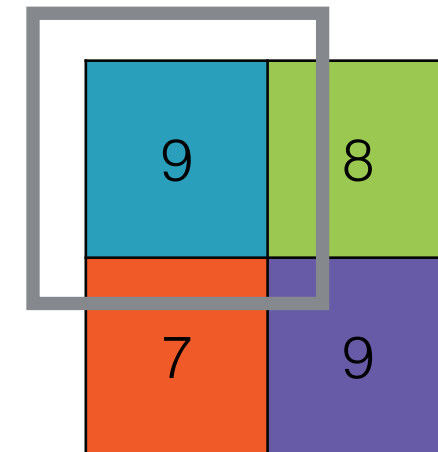


# Max Pooling



A 6x6 grid of numbers, color-coded by column: columns 1-3 are blue, 4-6 are green, 7-9 are orange, and 10-12 are purple. A 3x3 pooling kernel is highlighted by a gray border, covering the top-left 3x3 area of the grid.

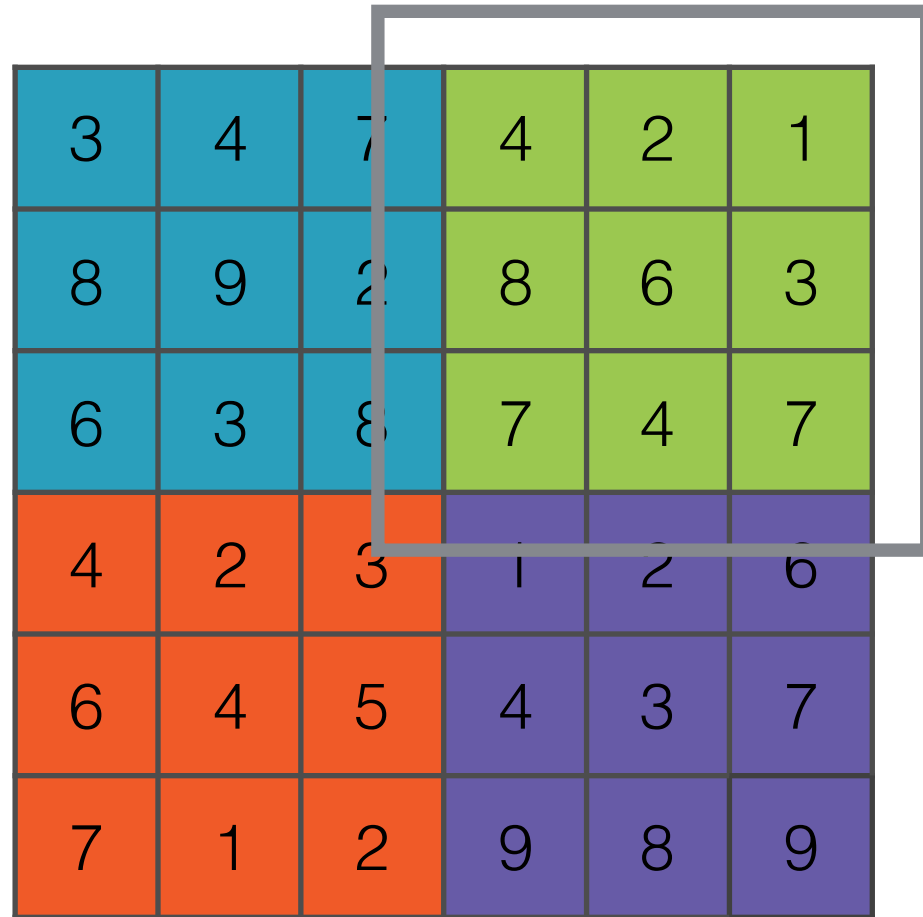
3	4	7	4	2	1
8	9	2	8	6	3
6	3	8	7	4	7
4	2	3	1	2	6
6	4	5	4	3	7
7	1	2	9	8	9



A 2x2 grid of numbers, color-coded by column: column 1 is blue, column 2 is green, column 3 is orange, and column 4 is purple. This grid represents the result of the max pooling operation applied to the 3x3 kernel from the previous grid.

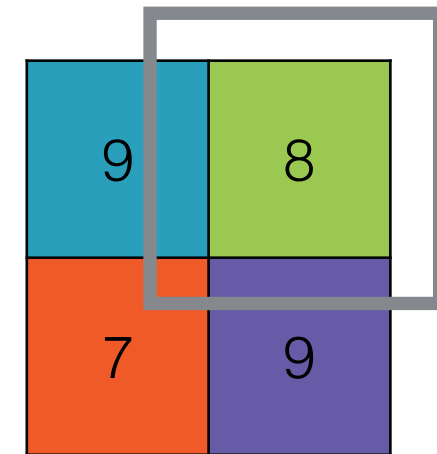
9	8
7	9

# Max Pooling



A 6x6 grid of numbers. The grid is divided into four colored regions: blue (top-left 3x3), green (top-right 3x3), orange (bottom-left 3x3), and purple (bottom-right 3x3). A gray border highlights a 3x3 max pooling kernel centered on the grid, covering the top-right 3x3 green region and the bottom-left 3x3 orange region.

3	4	7	4	2	1
8	9	2	8	6	3
6	3	8	7	4	7
4	2	3	1	2	6
6	4	5	4	3	7
7	1	2	9	8	9



A 2x2 grid of numbers representing the result of the max pooling operation. The grid is divided into four colored regions: blue (top-left), green (top-right), orange (bottom-left), and purple (bottom-right). A gray border highlights the 2x2 max pooling kernel, which is the top-right 2x2 green region and the bottom-left 2x2 orange region.

9	8
7	9

# Max Pooling

3	4	7	4	2	1
8	9	2	8	6	3
6	3	8	7	4	7
4	2	3	1	2	6
6	4	5	4	3	7
7	1	2	9	8	9

9	8
7	9

# Max Pooling

3	4	7	4	2	1
8	9	2	8	6	3
6	3	8	7	4	7
4	2	3	1	2	6
6	4	5	4	3	7
7	1	2	9	8	9

9	8
7	9

# Max Pooling

3	4	7	4	2	1
8	9	2	8	6	3
6	3	8	7	4	7
4	2	3	1	2	6
6	4	5	4	3	7
7	1	2	9	8	9

9	8
7	9

# Min Pooling

3	4	7	4	2	1
8	9	2	8	6	3
6	3	8	7	4	7
4	2	3	1	2	6
6	4	5	4	3	7
7	1	2	9	8	9

2	1
1	1

# Min Pooling

3	4	7	4	2	1
8	9	2	8	6	3
6	3	8	7	4	7
4	2	3	1	2	6
6	4	5	4	3	7
7	1	2	9	8	9


2	1
1	1

# Average Pooling

3	4	7	4	2	1
8	9	2	8	6	3
6	3	8	7	4	7
4	2	3	1	2	6
6	4	5	4	3	7
7	1	2	9	8	9

5.5	5.66
5.5	5.4





# Pooling Layers

**Pooling is widely used in Convolutional Neural Networks (CNNs)**

**Allow extraction of features in a location invariant manner**

Demo

**Block views on image arrays**

Demo

**Contour detection with the marching  
squares algorithm**

# Contour Line

A curve along connecting pixels with the same color;  
more generally a curve along which a function has the  
same value



# Contours

## **Alternative names for contour lines**

- equipotential curves
- indifference curves
- isopleths, isoclines

# Common Applications



## **Cartography (mapping)**

- connect points at equal elevation above sea level

## **Economics**

- connect points with equal utility (indifference curves)

# Two Types of Contours



**Isolines: Lines exactly follow single data value**

- the common data value is called the isovalue

**Isobands: Areas between isolines are filled in**

# The Marching Squares algorithm

An easy-to-implement, embarrassingly parallel algorithm that generates contours for a two-dimensional (rectangular) array



# Embarrassingly Parallel Algorithm

Where little or no effort is needed to separate the problem into a number of parallel tasks, usually because there is little or no dependency or need for communication between those parallel tasks


Demo

**Finding the convex hull of an image**

# Convex Hull

Given a  $X$ , a set of points in a plane, the convex hull of  $X$  is the smallest convex polygon that contains all points in  $X$

# Convex Hull



# Applications

**Widely used in image processing**

**Pattern recognition**

**Statistics**

**Many other disciplines**

# Edge Detection

---


# Edge

Curve connecting points where intensity (brightness) of pixels changes abruptly

# Edge Detection

Techniques that detect edges in images






# Applications

**Feature detection**

**Feature extraction**


**Widely used in computer vision**

# Types of Edge Detection



**Search-based**  
**Zero-crossing**

# Zero-crossing Edge Detection




**Zero-crossing: point along curve where value flips sign**

- Positive to negative
- Negative to positive

**In continuous curve, this occurs at value of zero**

# Search-based Edge Detection



**Find some measure of “edge strength”**

**First-order derivative suffices**

**Search for direction of maxima**

- Single maxima: edge
- Multiple maxima: corner

# Intuition behind Search-based Edge Detection

Compute  
derivatives in x and  
y direction

Find gradient  
magnitude

Apply threshold to  
gradient  
magnitude


# Intuition behind Search-based Edge Detection

Compute  
derivatives in x and  
y direction


Find gradient  
magnitude

Apply threshold to  
gradient  
magnitude

# Intuition behind Search-based Edge Detection




# Intuition behind Search-based Edge Detection





# Intuition behind Search-based Edge Detection



# Intuition behind Search-based Edge Detection

Compute  
derivatives in x and  
y direction

Find gradient  
magnitude

Apply threshold to  
gradient  
magnitude

# Intuition behind Search-based Edge Detection

Compute  
derivatives in x and  
y direction

Find gradient  
magnitude

**Apply threshold to  
gradient  
magnitude**

**Apply threshold to  
gradient  
magnitude**

**Pair of 2x2 convolution kernels**

**Designed to respond maximally to edges  
running at 45° to the pixel grid**

**Simple and quick operation**

# Convolution

In this context, a sliding window function applied to a matrix

# Representing Images

0	0	0	0	0	0
0.2	0.8	0	0.3	0.6	0
0.2	0.9	0	0.3	0.8	0
0.3	0.8	0.7	0.8	0.9	0
0	0	0	0.2	0.8	0
0	0	0	0.2	0.2	0

**Matrix**

1	0	1
0	1	0
1	0	1

**Kernel**

# Representing Images

0	0	0	0	0	0
0.2	0.8	0	0.3	0.6	0
0.2	0.9	0	0.3	0.8	0
0.3	0.8	0.7	0.8	0.9	0
0	0	0	0.2	0.8	0
0	0	0	0.2	0.2	0

1	0	1
0	1	0
1	0	1

**The kernel has weights which are tuned to  
extract edges**

# Representing Images

0	0	0	0	0	0
0.2	0.8	0	0.3	0.6	0
0.2	0.9	0	0.3	0.8	0
0.3	0.8	0.7	0.8	0.9	0
0	0	0	0.2	0.8	0
0	0	0	0.2	0.2	0

1	0	1
0	1	0
1	0	1

The weights for the Roberts and Sobel edge detection are different



# Convolution

0	0	0	0	0	0
0.2	0.8	0	0.3	0.6	0
0.2	0.9	0	0.3	0.8	0
0.3	0.8	0.7	0.8	0.9	0
0	0	0	0.2	0.8	0
0	0	0	0.2	0.2	0

Matrix



x1	x0	x1
x0	x1	x0
x1	x0	x1

1	1.2	1.1	0.9
1.9	2.7	2.5	1.9
1.0	2.1	2.4	1.4
1.0	1.8	2.0	1.8

Convolution  
Result

# Convolution

0 <sub>x1</sub>	x0	0 <sub>x1</sub>	0	0	0
x0	0.8 <sub>x1</sub>	x0	0.3	0.6	0
0.2 <sub>x1</sub>	x0	0 <sub>x1</sub>	0.3	0.8	0
0.3	0.8	0.7	0.8	0.9	0
0	0	0	0.2	0.8	0
0	0	0	0.2	0.2	0

Matrix



1			

Convolution  
Result

# Convolution

0	0 <sub>x1</sub>	x0	0 <sub>x1</sub>	0	0
0.2	x0	0 <sub>x1</sub>	x0	0.6	0
0.2	0.9 <sub>x1</sub>	x0	0.3 <sub>x1</sub>	0.8	0
0.3	0.8	0.7	0.8	0.9	0
0	0	0	0.2	0.8	0
0	0	0	0.2	0.2	0

Matrix



1	1.2		

Convolution  
Result

# Convolution

0	0	0 <sub>x1</sub>	x0	0 <sub>x1</sub>	0
0.2	0.8	x0	0.3 <sub>x1</sub>	x0	0
0.2	0.9	0 <sub>x1</sub>	x0	0.8 <sub>x1</sub>	0
0.3	0.8	0.7	0.8	0.9	0
0	0	0	0.2	0.8	0
0	0	0	0.2	0.2	0

Matrix



1	1.2	1.1	

Convolution  
Result

# Convolution

0	0	0	0 <sub>x1</sub>	x0	0 <sub>x1</sub>
0.2	0.8	0	x0	0.6 <sub>x1</sub>	x0
0.2	0.9	0	0.3 <sub>x1</sub>	x0	0 <sub>x1</sub>
0.3	0.8	0.7	0.8	0.9	0
0	0	0	0.2	0.8	0
0	0	0	0.2	0.2	0

Matrix



1	1.2	1.1	0.9

Convolution  
Result

# Convolution

0	0	0	0	0	0
0.2 <sub>x1</sub>	x0	0 <sub>x1</sub>	0.3	0.6	0
x0	0.9 <sub>x1</sub>	x0	0.3	0.8	0
0.3 <sub>x1</sub>	x0	0.7 <sub>x1</sub>	0.8	0.9	0
0	0	0	0.2	0.8	0
0	0	0	0.2	0.2	0

Matrix



1	1.2	1.1	0.9
1.9			

Convolution  
Result

# Convolution

0	0	0	0	0	0
0.2	0.8 <sub>x1</sub>	x0	0.3 <sub>x1</sub>	0.6	0
0.2	x0	0 <sub>x1</sub>	x0	0.8	0
0.3	0.8 <sub>x1</sub>	x0	0.8 <sub>x1</sub>	0.9	0
0	0	0	0.2	0.8	0
0	0	0	0.2	0.2	0

Matrix



1	1.2	1.1	0.9
1.9	2.7		

Convolution  
Result

# Convolution

0	0	0	0	0	0
0.2	0.8	0 <sub>x1</sub>	<b>x0</b>	0.6 <sub>x1</sub>	0
0.2	0.9	<b>x0</b>	0.3 <sub>x1</sub>	<b>x0</b>	0
0.3	0.8	0.7 <sub>x1</sub>	<b>x0</b>	0.9 <sub>x1</sub>	0
0	0	0	0.2	0.8	0
0	0	0	0.2	0.2	0

Matrix



1	1.2	1.1	0.9
1.9	2.7	2.5	

Convolution  
Result



# Convolution

0	0	0	0	0	0
0.2	0.8	0	0.3 <sub>x1</sub>	x0	0 <sub>x1</sub>
0.2	0.9	0	x0	0.8 <sub>x1</sub>	x0
0.3	0.8	0.7	0.8 <sub>x1</sub>	x0	0 <sub>x1</sub>
0	0	0	0.2	0.8	0
0	0	0	0.2	0.2	0

Matrix



1	1.2	1.1	0.9
1.9	2.7	2.5	1.9

Convolution  
Result

# Convolution

0	0	0	0	0	0
0.2	0.8	0	0.3	0.6	0
0.2 <sub>x1</sub>	x0	0 <sub>x1</sub>	0.3	0.8	0
x0	0.8 <sub>x1</sub>	x0	0.8	0.9	0
0 <sub>x1</sub>	x0	0 <sub>x1</sub>	0.2	0.8	0
0	0	0	0.2	0.2	0

Matrix



1	1.2	1.1	0.9
1.9	2.7	2.5	1.9
1.0			

Convolution  
Result

# Convolution

0	0	0	0	0	0
0.2	0.8	0	0.3	0.6	0
0.2	0.9 <sub>x1</sub>	x0	0.3 <sub>x1</sub>	0.8	0
0.3	x0	0.7 <sub>x1</sub>	x0	0.9	0
0	0 <sub>x1</sub>	x0	0.2 <sub>x1</sub>	0.8	0
0	0	0	0.2	0.2	0

Matrix



1	1.2	1.1	0.9
1.9	2.7	2.5	1.9
1.0	2.1		

Convolution  
Result

# Convolution

0	0	0	0	0	0
0.2	0.8	0	0.3	0.6	0
0.2	0.9	0 <sub>x1</sub>	<b>x0</b>	0.8 <sub>x1</sub>	0
0.3	0.8	<b>x0</b>	0.8 <sub>x1</sub>	<b>x0</b>	0
0	0	0 <sub>x1</sub>	<b>x0</b>	0.8 <sub>x1</sub>	0
0	0	0	0.2	0.2	0

Matrix



1	1.2	1.1	0.9
1.9	2.7	2.5	1.9
1.0	2.1	2.4	

Convolution  
Result

# Convolution

0	0	0	0	0	0
0.2	0.8	0	0.3	0.6	0
0.2	0.9	0	0.3 <sub>x1</sub>	x0	0 <sub>x1</sub>
0.3	0.8	0.7	x0	0.9 <sub>x1</sub>	x0
0	0	0	0.2 <sub>x1</sub>	x0	0 <sub>x1</sub>
0	0	0	0.2	0.2	0

Matrix



1	1.2	1.1	0.9
1.9	2.7	2.5	1.9
1.0	2.1	2.4	1.4

Convolution  
Result

# Convolution

0	0	0	0	0	0
0.2	0.8	0	0.3	0.6	0
0.2	0.9	0	0.3	0.8	0
0.3 <sub>x1</sub>	<b>x0</b>	0.7 <sub>x1</sub>	0.8	0.9	0
<b>x0</b>	0 <sub>x1</sub>	<b>x0</b>	0.2	0.8	0
0 <sub>x1</sub>	<b>x0</b>	0 <sub>x1</sub>	0.2	0.2	0

Matrix



1	1.2	1.1	0.9
1.9	2.7	2.5	1.9
1.0	2.1	2.4	1.4
1.0			

Convolution  
Result

# Convolution

0	0	0	0	0	0
0.2	0.8	0	0.3	0.6	0
0.2	0.9	0	0.3	0.8	0
0.3	0.8 <sub>x1</sub>	<b>x0</b>	0.8 <sub>x1</sub>	0.9	0
0	<b>x0</b>	0 <sub>x1</sub>	<b>x0</b>	0.8	0
0	0 <sub>x1</sub>	<b>x0</b>	0.2 <sub>x1</sub>	0.2	0

Matrix



1	1.2	1.1	0.9
1.9	2.7	2.5	1.9
1.0	2.1	2.4	1.4
1.0	1.8		

Convolution  
Result

# Convolution

0	0	0	0	0	0
0.2	0.8	0	0.3	0.6	0
0.2	0.9	0	0.3	0.8	0
0.3	0.8	0.7 <sub>x1</sub>	x0	0.9 <sub>x1</sub>	0
0	0	x0	0.2 <sub>x1</sub>	x0	0
0	0	0 <sub>x1</sub>	x0	0.2 <sub>x1</sub>	0

Matrix



1	1.2	1.1	0.9
1.9	2.7	2.5	1.9
1.0	2.1	2.4	1.4
1.0	1.8	2.0	

Convolution  
Result



# Convolution

0	0	0	0	0	0
0.2	0.8	0	0.3	0.6	0
0.2	0.9	0	0.3	0.8	0
0.3	0.8	0.7	0.8 <sub>x1</sub>	<b>x0</b>	0 <sub>x1</sub>
0	0	0	<b>x0</b>	0.8 <sub>x1</sub>	<b>x0</b>
0	0	0	0.2 <sub>x1</sub>	<b>x0</b>	0 <sub>x1</sub>

Matrix



1	1.2	1.1	0.9
1.9	2.7	2.5	1.9
1.0	2.1	2.4	1.4
1.0	1.8	2.0	1.8

Convolution  
Result

# Demo

## **Performing edge detection using**

- Roberts cross detection
- Sobel edge detection

# Canny Edge Detection

---

# Canny Edge Detection

Multi-stage algorithm, named after its inventor, that detects useful structural information from images with greatly reduced data processing

Enhances the signal to noise ratio  
thus producing better edge  
detection with noisy images

# Canny Edge Detection

## Smoothen Image

Remove noise with Gaussian filter

## Suppress Spurious Responses

Set all gradient values except maxima to 0

## Edge Tracking by Hysteresis

Eliminate all weak edges not connected to strong edges

## Compute Intensity Gradients


Edge gradient and direction as before

## Apply Double Threshold

Filter out edge pixels with weak gradients; preserve edge pixels with strong gradients


# Gaussian Filter for Smoothing

**Start with a  
set of points  
in space**



# Gaussian Filter for Smoothing


**Define a  
neighborhood  
for each point**






# Gaussian (RBF) Kernel

**Probability-weighted  
sum of points**



**What probability  
distribution?**

# Gaussian (RBF) Kernel




**Gaussian probability distribution**

**Defined by**

- mean  $\mu$
- standard deviation  $\sigma$

Mean = Center point




# Gaussian Filter

**Mean  $\mu$  = center point**

**Standard deviation  $\sigma$  ~ bandwidth**

Mean = Center point



## Gaussian (RBF) Kernel

**Larger value of sigma:** less sensitive to noise, detects only large edges

**Smaller value of sigma:** more sensitive to noise, detects even fine edges

**Need to strike a fine balance**

**Sigma influences quality of feature extraction**

Demo

**Performing Canny edge detection**

# Summary

**Working with images is increasingly important**

**scikit-image is an image processing toolkit**

**Performed image manipulation, contour and convex hull determination**

**Edge detection using Sobel, Roberts and Canny edge detectors**