

Employing Ensemble Methods with scikit-learn

UNDERSTANDING ENSEMBLE LEARNING TECHNIQUES



Janani Ravi

CO-FOUNDER, LOONYCORN

www.loonycorn.com

Overview

Ensemble learning to improve robustness and reduce overfitting

Different kinds of ensemble learning techniques

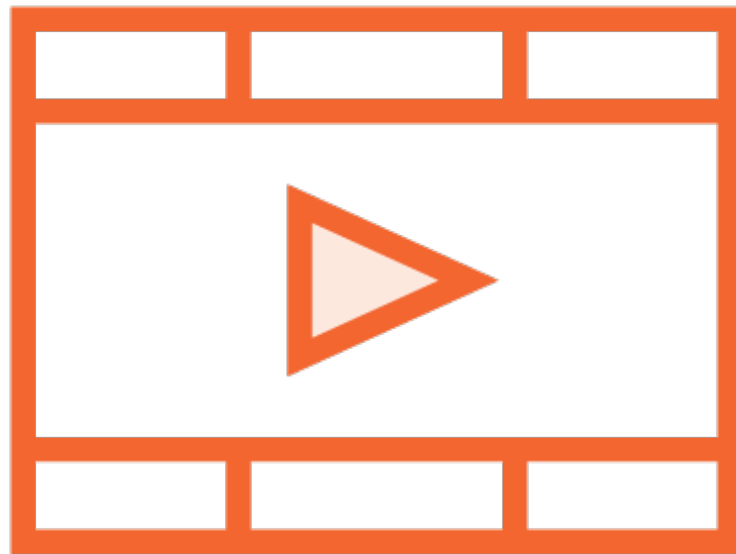
Averaging, boosting, voting, stacking

Built-in support for ensemble learning in scikit-learn

Implementing hard and soft voting in scikit-learn

Prerequisites and Course Outline

Prerequisites

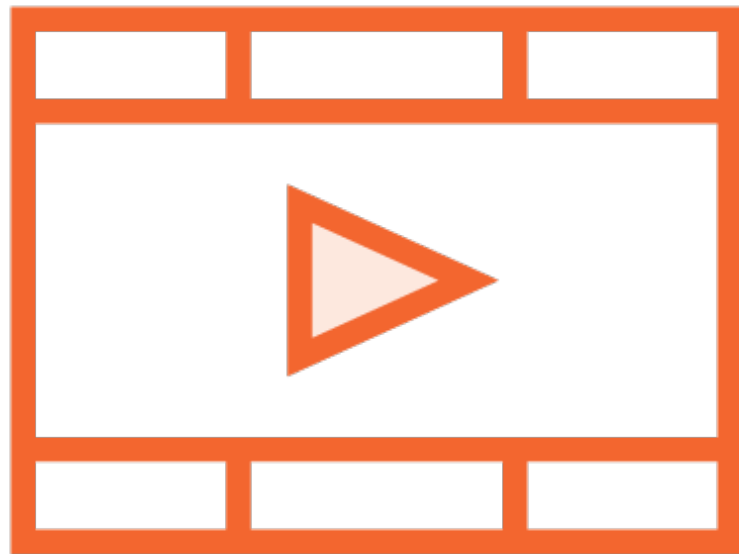


Comfortable with Python programming

Prior ML exposure recommended

Building simple classification and regression models using scikit-learn

Prerequisite Courses



Building Your First scikit-learn Solution

**Building Classification Models with
scikit-learn**

**Building Regression Models with
scikit-learn**

Course Outline



Introducing ensemble learning

**Ensemble learning using averaging -
bagging and pasting models**

**Ensemble learning using boosting -
adaptive and gradient boosting**

Ensemble learning using stacking

Quick Overview of Ensemble Learning

Ensemble Learning

Machine learning technique in which several learners are combined to obtain a better performance than any of the learners individually.


Ensemble Learning

Machine learning technique in which **several learners are combined** to obtain a better performance than any of the learners individually.

Ensemble Learning

Machine learning technique in which several learners are combined to obtain a **better performance than any of the learners individually.**

Ensemble Learning



Important Questions in Ensemble Learning

**What kind of
individual learners
to use?**

**How should
individual learners
be trained?**

**How should
individual learners
be combined?**

Important Questions in Ensemble Learning

**What kind of
individual learners
to use?**

How should
individual learners
be trained?

How should
individual learners
be combined?

Choice of Individual Learners



Individual learners (models) could be of absolutely any type

Each learner should be as **different as possible from other learners**

Choice of Individual Learners



Decision trees are most often used

An ensemble of decision trees is a
Random Forest

Random forests make it easy to build
uncorrelated learners


Important Questions in Ensemble Learning

What kind of
individual learners
to use?

How should
individual learners
be trained?

How should
individual learners
be combined?

Training Individual Learners



If learners are **different**, each learner can be trained on the entire dataset

For **similar** learners:

- Each model is trained on **random samples** of training data
- Can also use **random set of features** to train different models

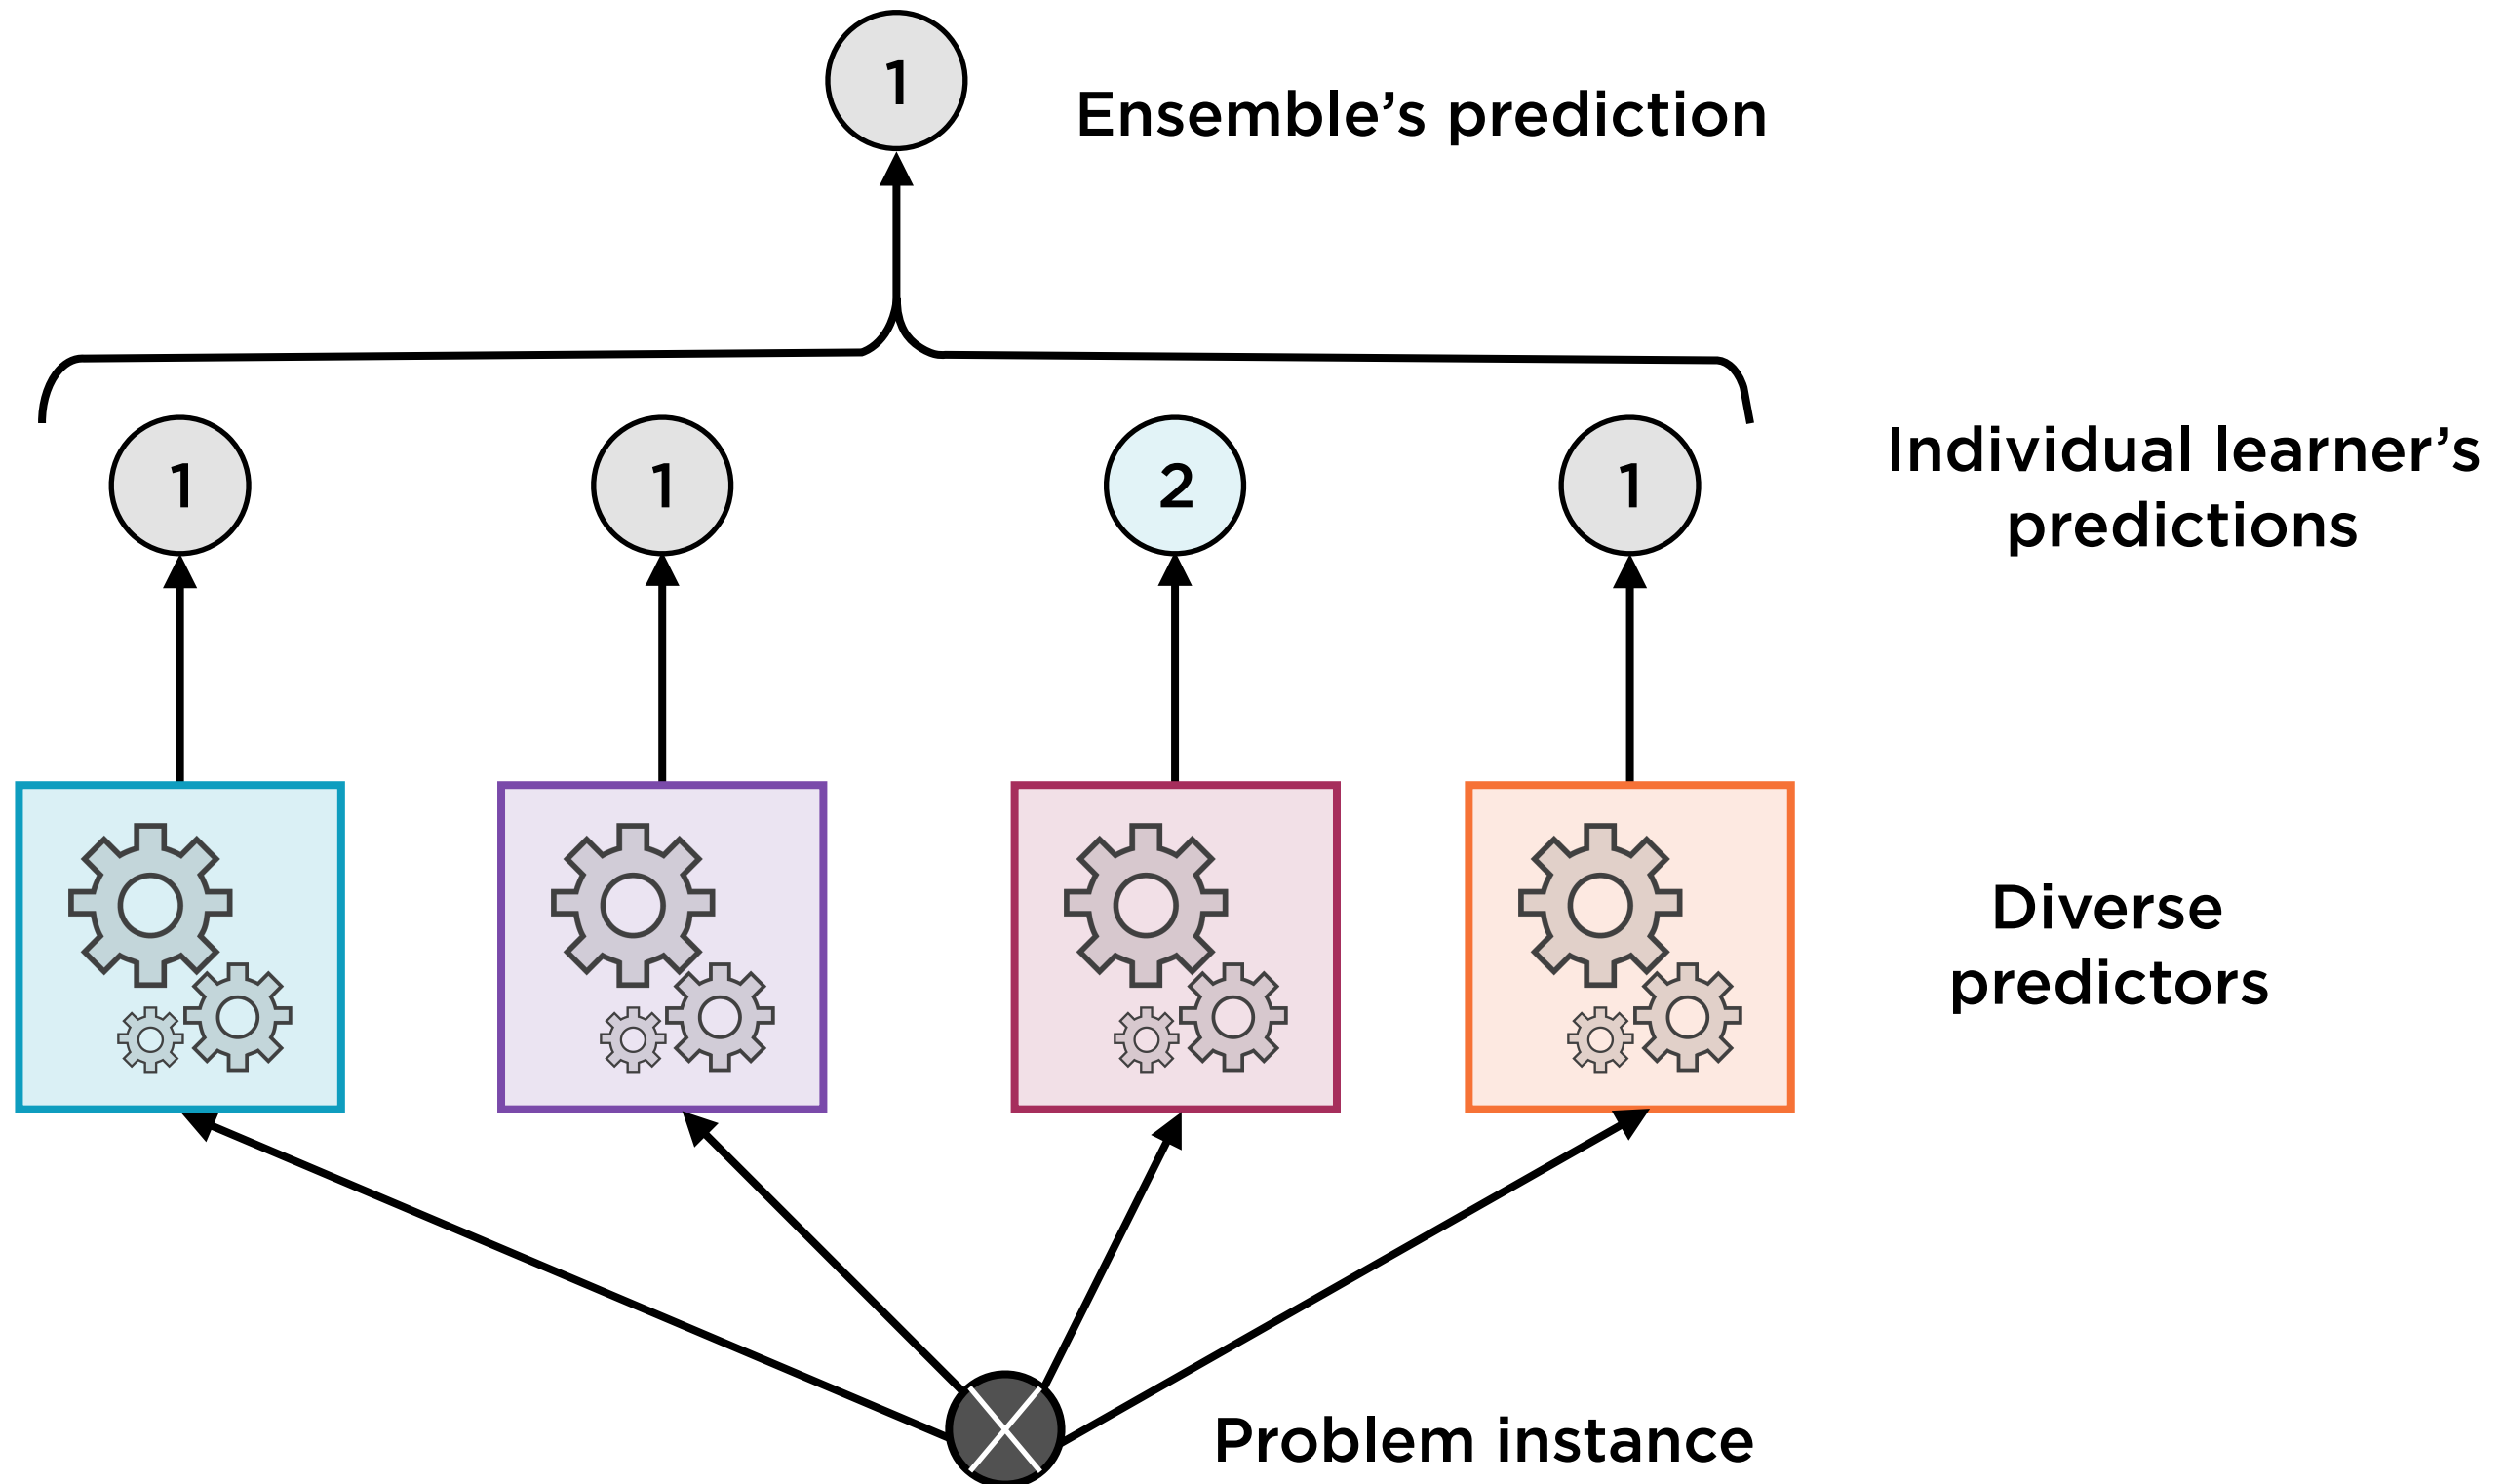
Important Questions in Ensemble Learning

What kind of
individual learners
to use?

How should
individual learners
be trained?

How should
individual learners
be combined?

Combining Classifier Predictions



Combining Individual Learners




Hard voting: Majority vote of individual learners (classification)

Soft voting: Probability-weighted average

Stacking: Train additional model to combine predictions from individual learners

Ensemble Learning Techniques

Averaging and Boosting



Averaging

**Train predictors in parallel and
average scores of individual
predictors**



Boosting


**Train predictors in sequence
where each predictor learns from
earlier mistakes**

Voting and Stacking



Voting

Majority vote of the individual predictors is the final prediction of the ensemble



Stacking

Fit a model on the individual predictions to get the final prediction of the ensemble

Voting



Get individual predictions from each learner

Each learner uses a different training algorithm

The different algorithms add diversity to the predictions

Voting




Hard Voting: Final output of the ensemble is the **majority** vote

Soft Voting: Final output of the ensemble is the category with the **highest probability** score

- Need to be able to **aggregate probability scores** for each output category

Averaging




Train multiple learners in **parallel**

Get individual predictions from each learner

Final prediction of the ensemble is an **average** of individual predictions

Voting can be considered an averaging technique

Averaging



Usually use decision trees or random forests to build different models

Train model on different samples of training data

- Bagging: Sample data **with replacement**
- Pasting: Sample data **without replacement**

Boosting



Train multiple learners **sequentially**

Each model learns from the mistakes made by previous models

Can tweak the learning rate or contribution of each model

Addition of a learner boosts the accuracy of the model


Boosting



Adaptive Boosting: each model pays more attention to training instances the previous model got wrong

Gradient Boosting: each model in sequence fits on residual errors of the previous model

Stacking



Train diverse individual learners

Get predictions from individual predictors

Fit a model to make the final predictions of the ensemble

“Blender model” or “Meta-learner”

Decision Trees in Ensemble Learning

Important Questions in Ensemble Learning

**What kind of
individual learners
to use?**

How should
individual learners
be trained?

How should
individual learners
be combined?

Choice of Individual Learners



Individual learners (models) could be of absolutely any type

Could combine:

- Neural networks
- Support Vector Machines
- Naive Bayes Classifiers
- Decision Trees
- Random Forests (group of decision trees)

Ensemble Learning



Individual learners should be as **different** as possible

For most techniques, hard to generate large number of very different models

To the rescue: Decision trees and random forests

Decision Trees are the most
common building blocks for
Ensemble Learning

Decision Trees

ML models that construct trees based on threshold values of x -variables. Differ from rule-based trees because thresholds are determined by training.

Random Forest

An ensemble (collection) of decision trees, in which individual trees are trained on different random subsets of training data.


Important Questions in Ensemble Learning

What kind of
individual learners
to use?

How should
individual learners
be trained?

How should
individual learners
be combined?


Training Individual Models



If individual learners are not decision trees

- Then each has independent, full training process

Training Random Forests



Most common form of ensemble learning uses random forests

Individual learners are decision trees

Each tree is iteratively trained on randomly sampled subset

Training Random Forests



Will return to this in a later module

Decision Trees

Jockey or Basketball Player?



Jockeys

Tend to be light to meet horse carrying limits



Basketball Players

Tend to be tall, strong and heavy

Jockey or Basketball Player?



Intuitively know

Jockeys tend to be light

And not very tall


Basketball players tend to be tall

And also quite heavy




Decision trees set up a tree structure on training data which helps make **decisions** based on **rules**


Fit Knowledge into Rules




Decision Based on Weight




Decision Based on Weight



Fit Knowledge into Rules




Fit Knowledge into Rules



Decision Tree

Fit knowledge
into rules


Each rule involves
a threshold



Decision Tree

Order of decision
variables matters


Rules and order
found using ML



Splitting a Decision Tree

Tree selects the **best** feature


And finds the best **threshold** for the feature



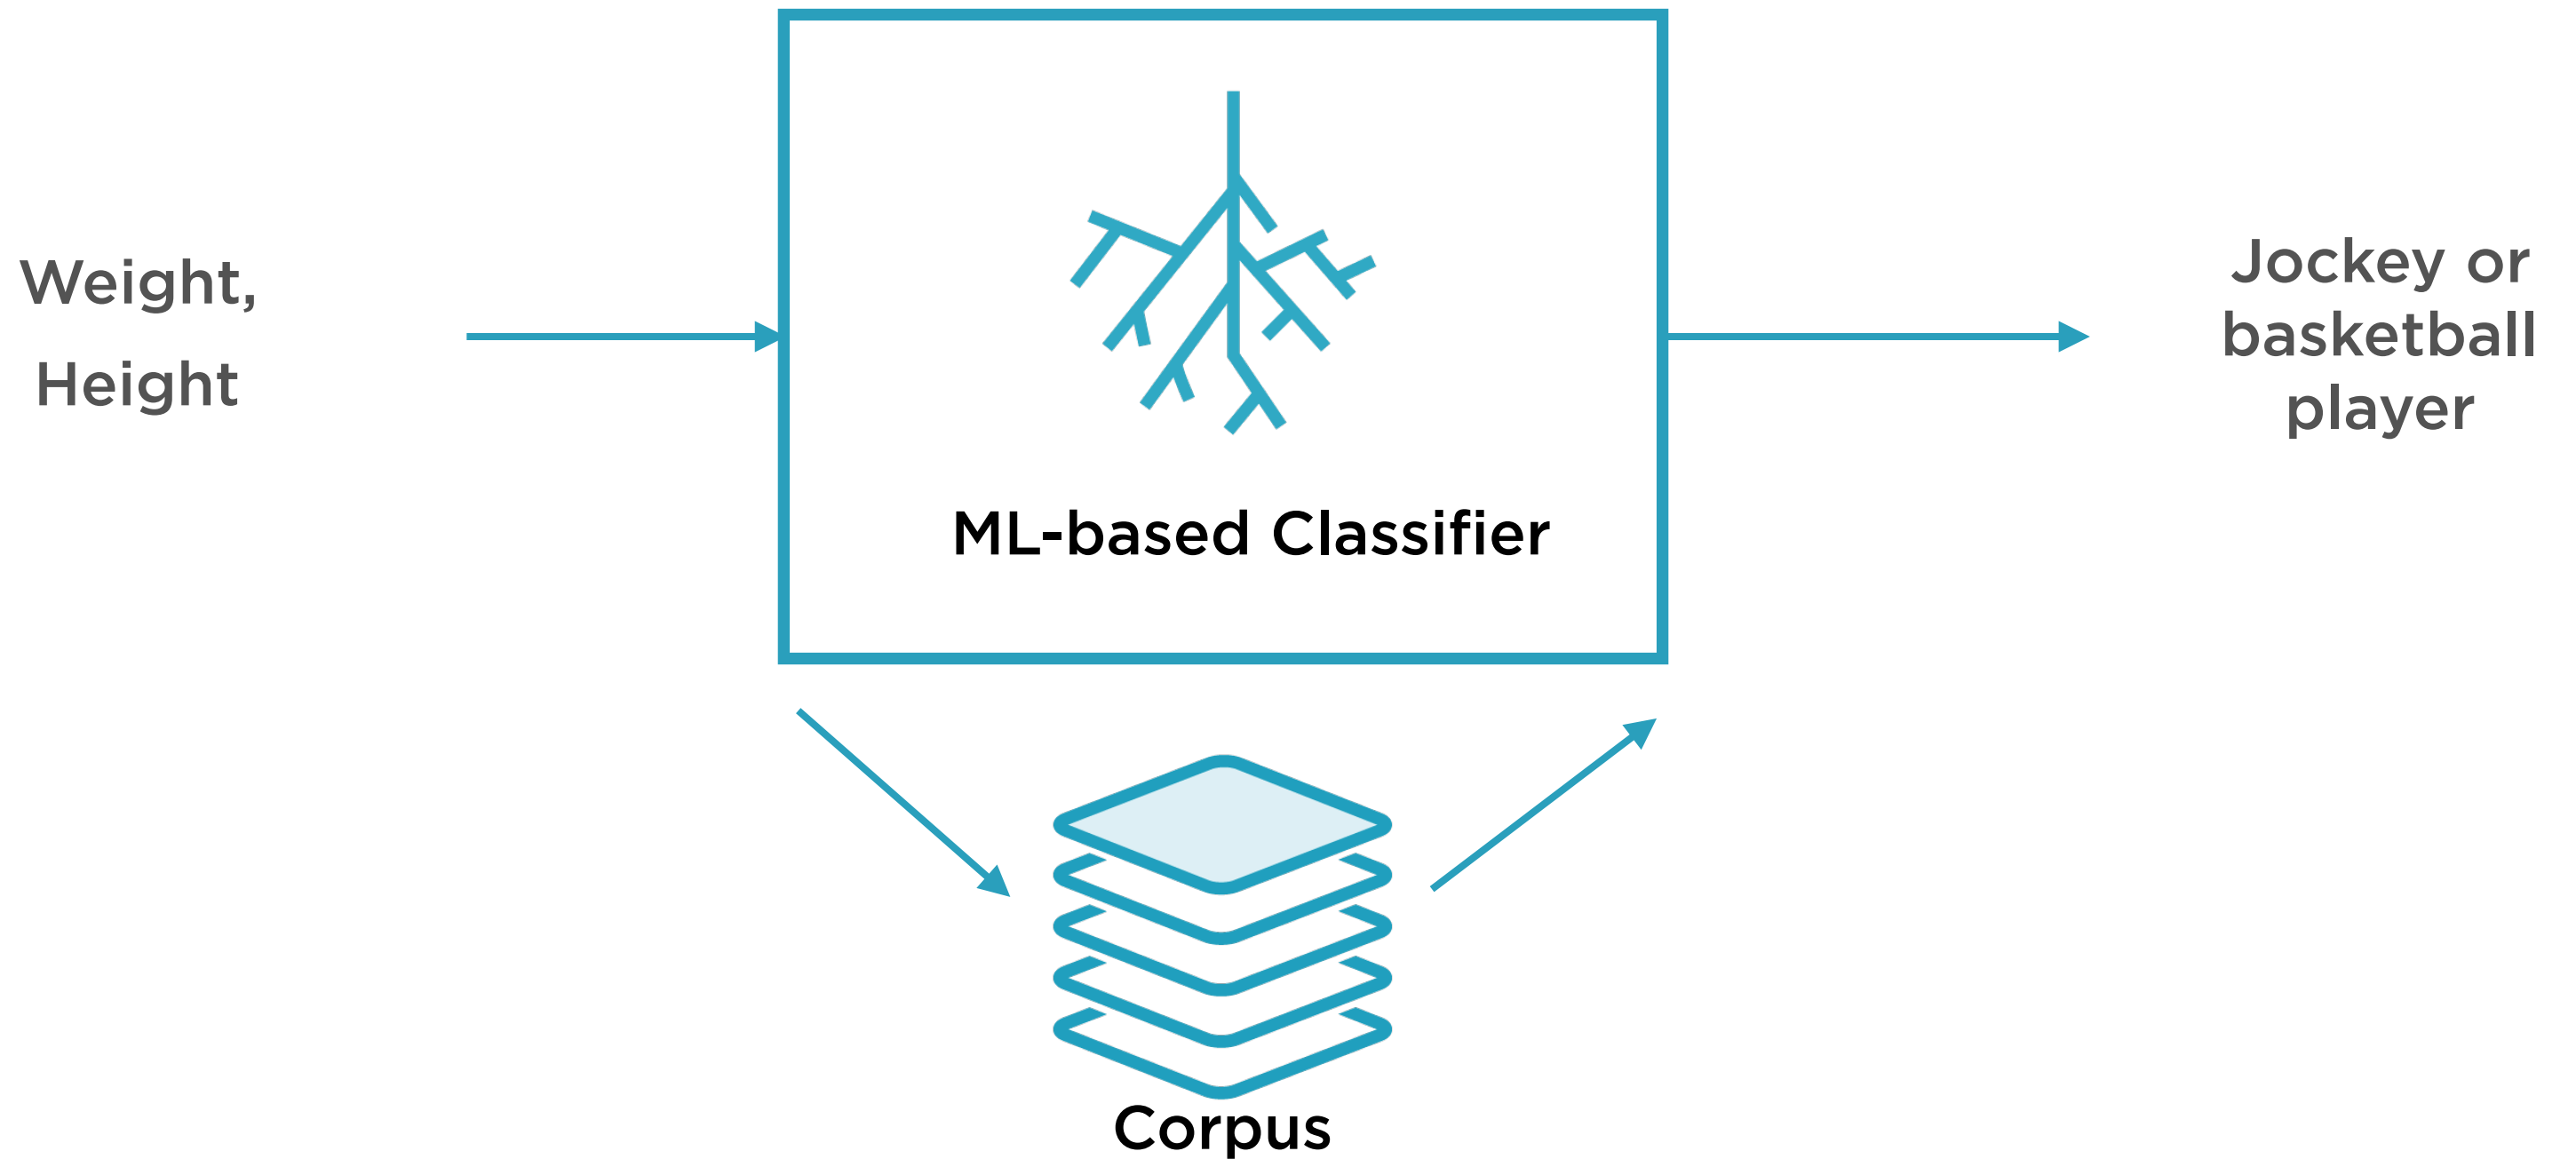
Decision Tree

“CART”


Classification And
Regression Tree



Decision Trees for Classification



Decision Trees for Classification




Traverse tree to find right node

Return **most frequent label** of all training data points in that node


Ensemble Learning to Mitigate Overfitting

Connecting the Dots




Challenge: Fit the “best” curve through these points

Good Fit?




A curve has a “good fit” if the distances of points from the curve are small

Connecting the Dots




We could draw a pretty complex curve

Connecting the Dots




We can even make it pass through every single point

Connecting the Dots




But given a new set of points, this curve might perform quite poorly

Connecting the Dots




The original points were “training data”, the new points are “test data”

Overfitting




Great performance in training, poor performance in real usage

Connecting the Dots



A simple straight line performs worse in training, but better with test data

Overfitting




Model has memorized the training data

Low training error

Does not work well in the real world

High test error

Cause of Overfitting




Sub-optimal choice in the **bias-variance** trade-off

An overfitted model has:

- high variance error
- low bias error




Bias



Low bias

Few assumptions about the
underlying data




High bias

More assumptions about the
underlying data




Bias



Model too complex

Training data all-important, model
parameter counts for little




Model too simple

Model parameter all-important,
training data counts for little




Variance



High variance

The model changes significantly
when training data changes




Low variance

The model doesn't change much
when the training data changes




Variance



Model too complex


Model varies too much with changing
training data



Model too simple


Model not very sensitive to training
data

Bias-variance Trade-off



Model too complex

High variance error




Model too simple

High bias error

Preventing Overfitting



Regularization - Penalize complex models



Cross-validation - Distinct training and validation phases



Dropout (NNs only) - Intentionally turn off some neurons during training

Ensemble Learning an
important technique to
mitigate overfitting

Demo

Exploring the environment and tools

Demo

**Performing hard and soft voting using
the VotingClassifier**

Summary

Ensemble learning to improve robustness and reduce overfitting

Different kinds of ensemble learning techniques

Averaging, boosting, voting, stacking

Built-in support for ensemble learning in scikit-learn

Implementing hard and soft voting in scikit-learn