

Building Convolutional NN with Keras



Jerry Kurata

CONSULTANT


@jerrykur




... a convolutional neural network (CNN, or ConvNet) is a class of deep, feed-forward artificial NN that has successfully been applied to analyzing visual imagery.

Wikipedia





Curse of Dimensionality



More connections



More weights to train



Longer training


Curse of Dimensionality

Inputs

(8 MP)



Dense(1000)



8,000,000 pixels
X 1,000 neurons

8,000,000,000 weights
X 3 colors

24,000,000,000 weights





Translational Invariance



Convolutional Neural Network




Reduce weights to train




General object identification




Convolutional Neural Network




CNN Components



Convolution



Non-linearity
(ReLU)



Pooling

Cat 0.6
Dog 0.25
Tree 0.15

Classification



Convolution

Extracts features from image

Preserves feature spatial relationships

- Edges
- Composite elements (nose, eye)

Reduced computation



Filters

Values not fixed

Values are what we train

Improved through training

Trained on labeled images

Detect unique features that determine objects

CNN training faster since only filter weights trained

Weights shared across image



Key Convolution Hyperparameters

Kernel size

Number of filters

Stride

Padding



Kernel Size and Number of Filters

Kernel size -> related pixels


3X3, 5X5, 7X7

Number of Filters -> Feature
Detected

16, 32, 64, 96




kernel_size = (2,2)
filters = 5




Image

6 X 6 X 3



Kernel (Filters)

2 X 2 X 5



Feature Maps

5 X 5 X 5



Stride

Distance to move filter

Larger stride -> pixel
independence



Stride

Distance to move filter


Larger values faster

- Decrease size of feature map
- Reduces information passed to next layer

1 is common value




Stride = 3




Image

6 X 6 X 3



Kernel (Filters)

3 X 3 X 5



Feature Maps

2 X 2 X 5




Issues with Convolution

0 _{x1}	1 _{x0}	1 _{x1}	1	1	1
0 _{x0}	0 _{x0}	0 _{x0}	1	1	0
0 _{x1}	0 _{x0}	1 _{x1}	1	1	1
1	0	0	1	1	1
1	1	0	0	1	1
0	0	1	1	0	0

2	3	4	4
1	2	2	3
3	2	3	3
2	2	3	3


Reduction in spatial dimensions



Data at edges is used less




Padding = 'same'




Image

$6 \times 6 \times 3$
+ padding



Kernel (Filters)

$3 \times 3 \times 5$




Feature Maps

$6 \times 6 \times 5$



Non-Linear Activation Function



Lets NN handle non-linear

Added in two ways


- As a layer after convolution layer
- As parameter to convolution layer

ReLU most common non-linearity function

- $y = \max(0, x)$
- If $(x < 0)$ return 0 else return x

Prevents vanishing gradient

Convolutional Neural Network



Pooling Layer

3	5	7	2	3	3
8	1	6	1	1	4
5	6	9	3	2	9
3	2	9	5	7	9
5	4	7	4	8	9
2	5	8	3	3	1

8	7	4
6	9	

Reduce Dimensionality




Translational Invariance



Max Pooling

5	8	7	2	6	6
3	3	6	4	8	7
3	2	4	7	4	2
2	2	8	5	2	1
3	8	3	3	2	2
7	5	3	3	1	1

6 X 6 X 5



2 X 2

8	7	8
3	8	4
8	3	2

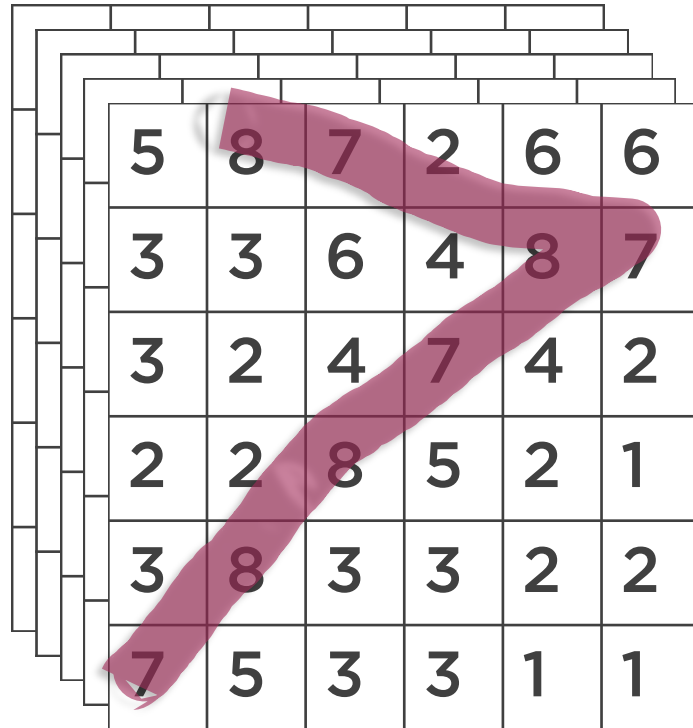
$$W_{\text{out}} = ((W_{\text{in}} - F)/S) + 1$$

$$H_{\text{out}} = ((H_{\text{in}} - F)/S) + 1$$

$$D_{\text{in}} = D_{\text{out}}$$




Max Pooling

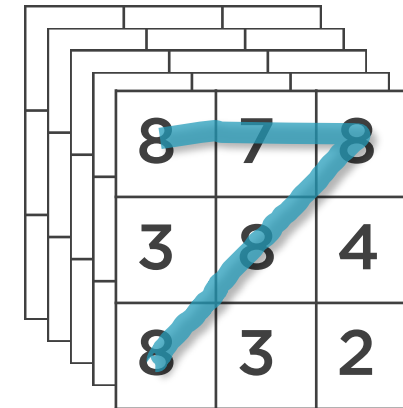


5	8	7	2	6	6
3	3	6	4	8	7
3	2	4	7	4	2
2	2	8	5	2	1
3	8	3	3	2	2
7	5	3	3	1	1

6 X 6 X 5



2 X 2



8	7	8
3	8	4
8	3	2

3 X 3 X 5


$$W_{\text{out}} = ((6 - 2)/2) + 1 = 3$$

$$H_{\text{out}} = ((6 - 2)/2) + 1 = 3$$

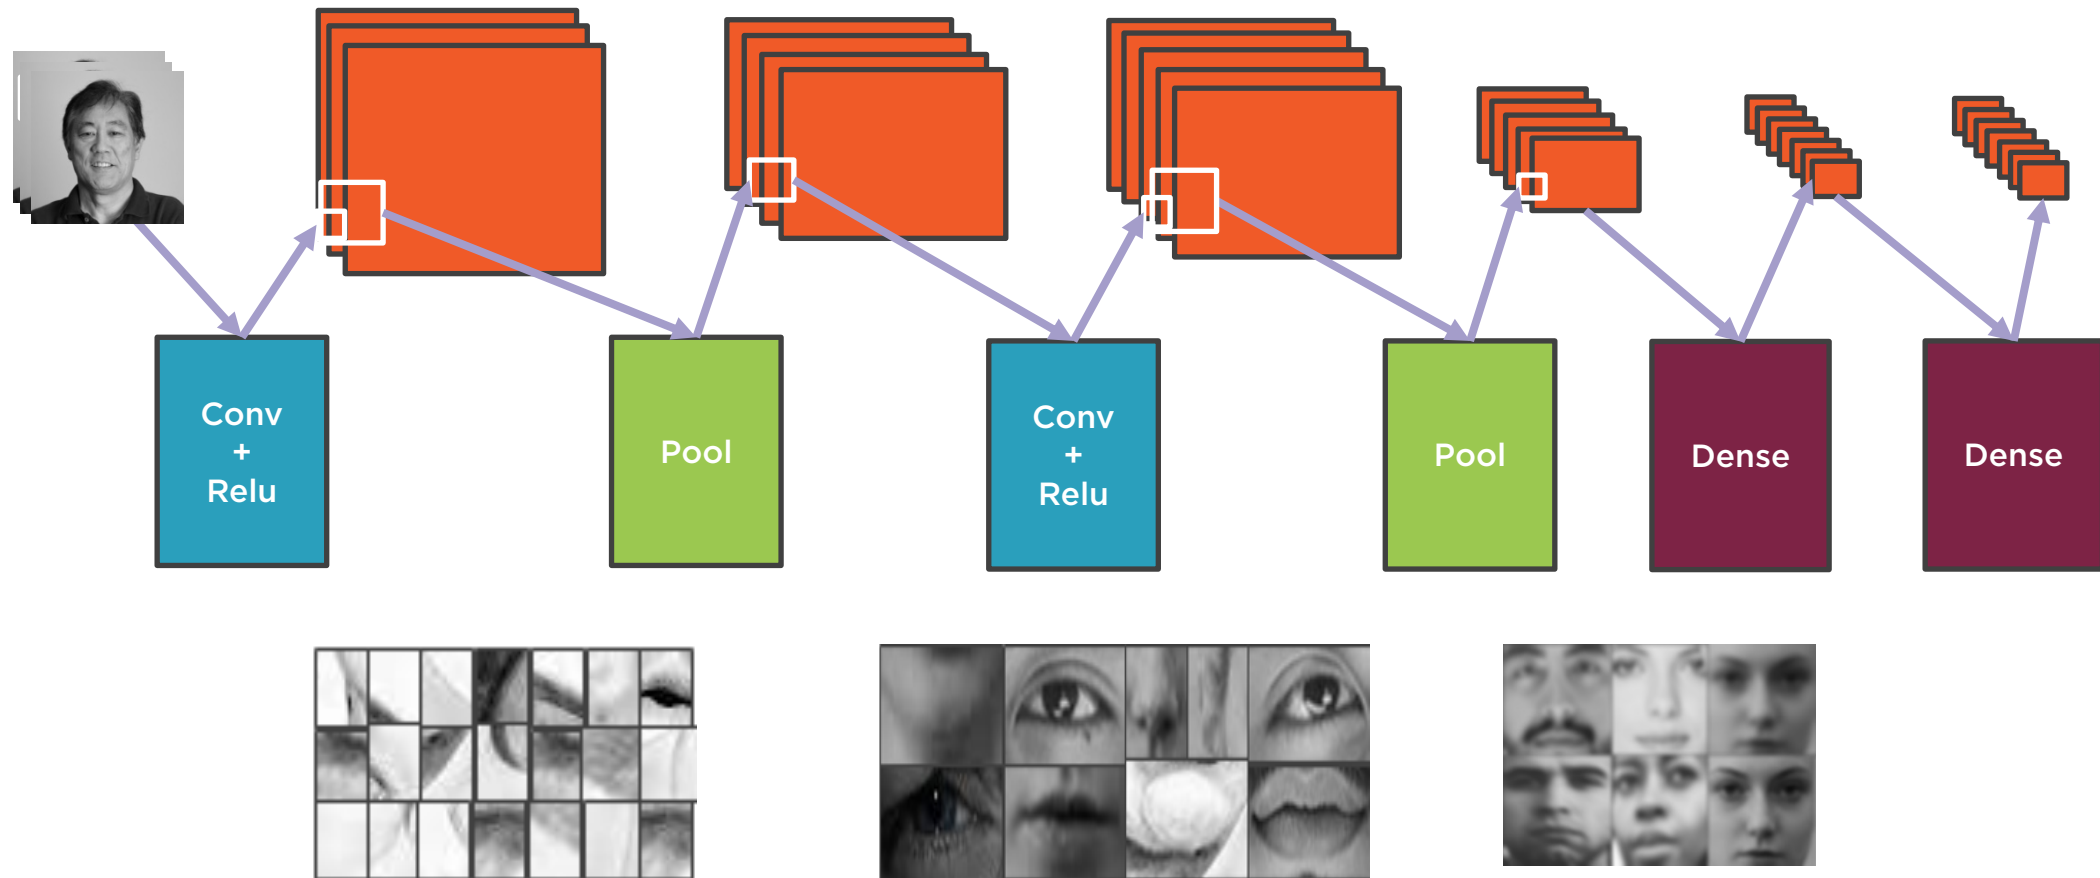
$$D_{\text{in}} = D_{\text{out}} = 5$$



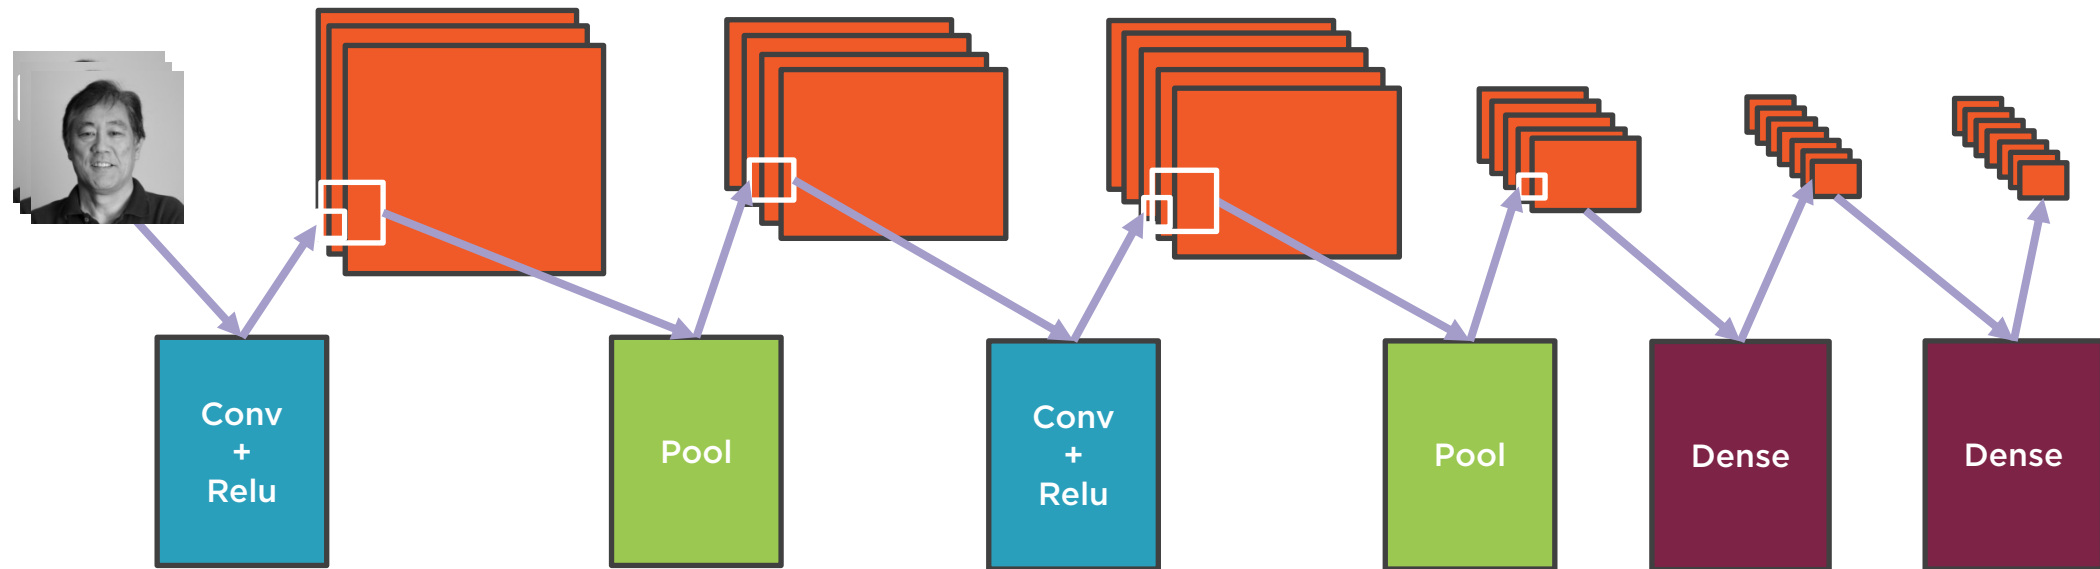
Convolutional Neural Network



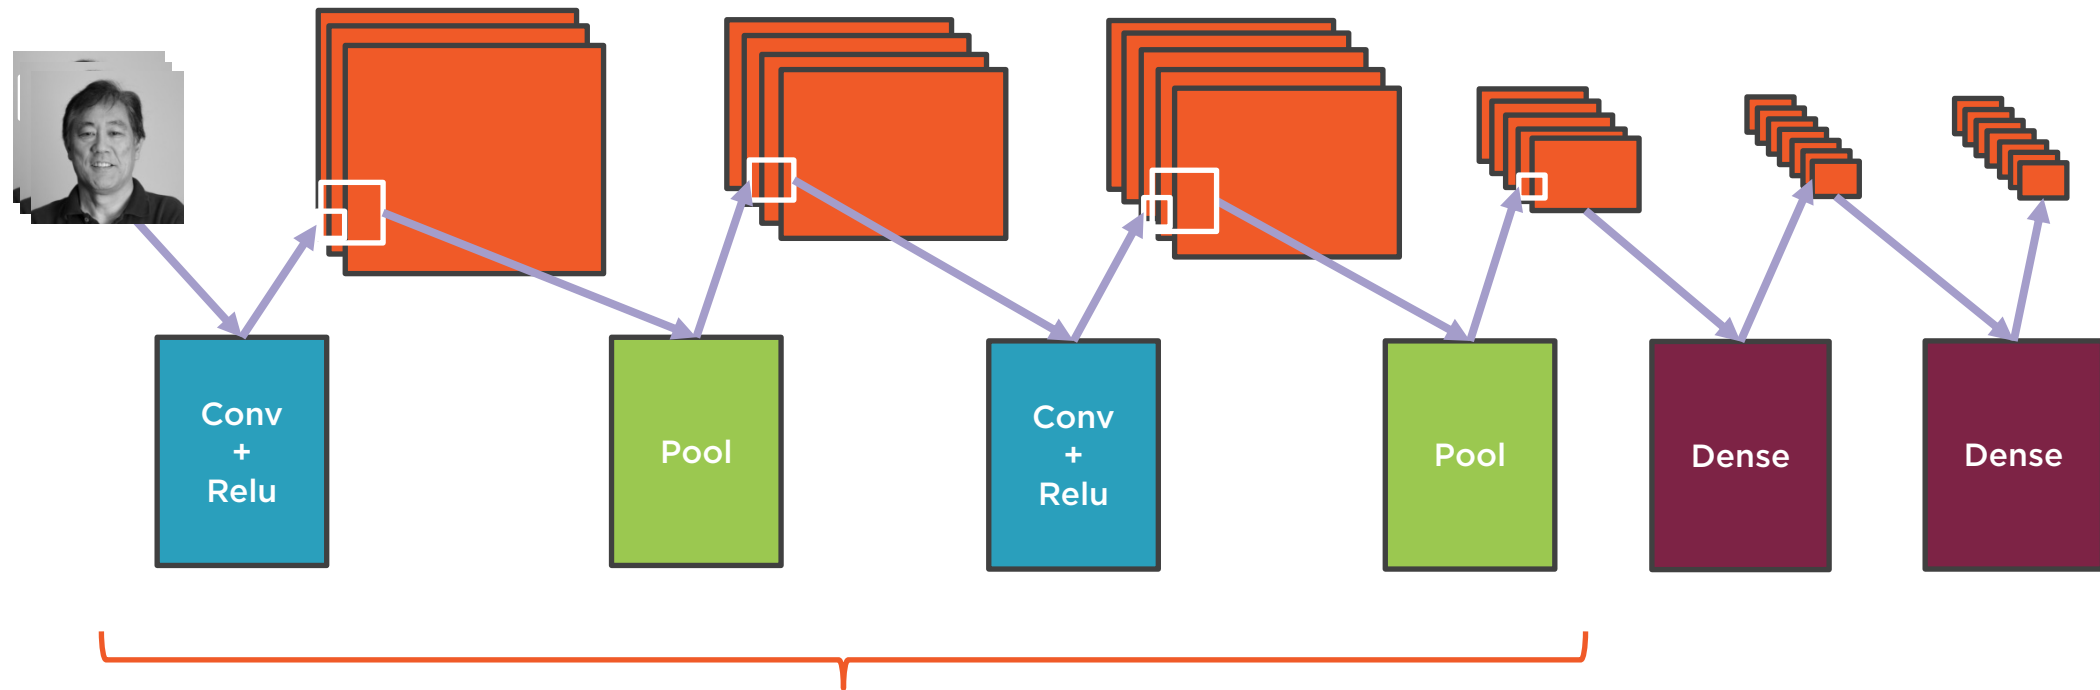
Convolutional Neural Network



Convolutional Neural Network



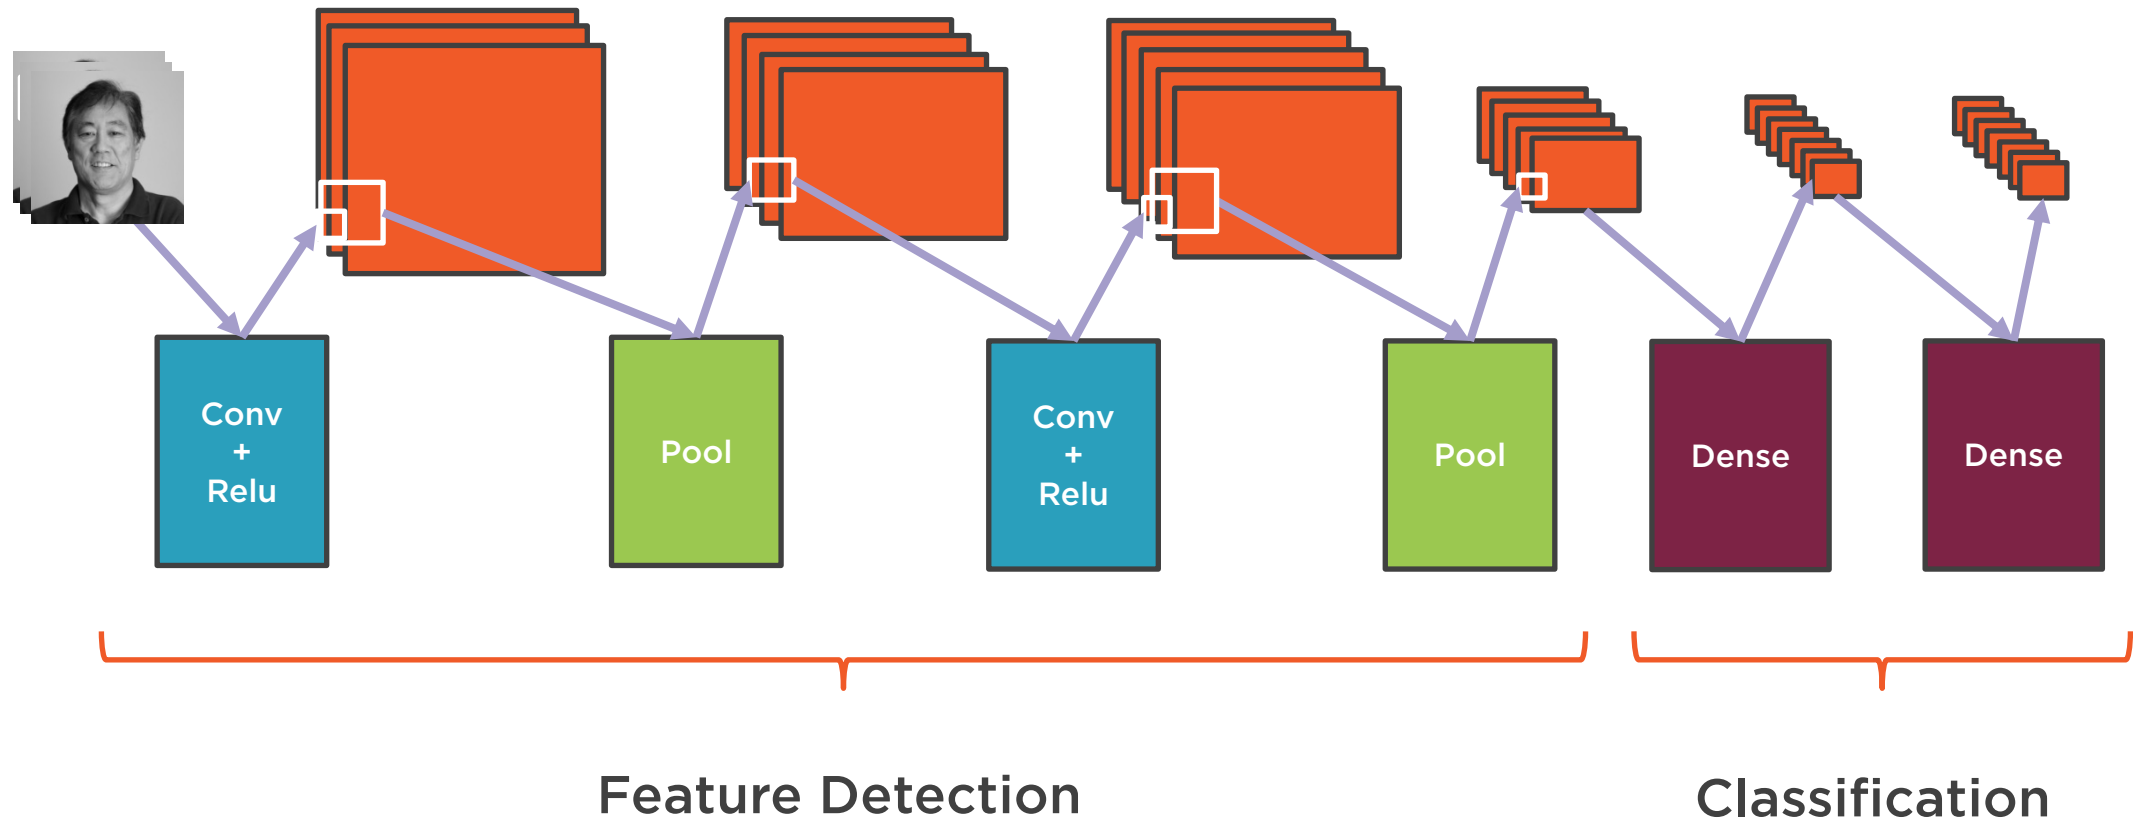
Convolutional Neural Network




Feature Detection



Convolutional Neural Network



Convolutional Neural Network





MNIST – Modified NIST

MNIST Handwritten digits

- 60,000 training images
- 10,000 testing images

28 X 28 X 1 (grayscale)


10 classes

- Digits 0 through 9

A basic CNN trained easily

Over 99% accuracy often achieved

Not a common task



Fashion-MNIST images

- 60,000 training
- 10,000 test

28 X 28 X 1 (grayscale)

10 classes

- (t-shirt, trouser, pullover, coat, sandal, skirt, sneaker, bag, and ankle boot)

Solvable by basic CNN

90% accuracy good, higher with tuning

Common task

Demo



Fashion MNIST Image Classification





Need powerful model


Don't want to build from scratch

Transfer Learning

- Leverage existing trained model
- Tailor to our problem

Fashion MNIST

Fashion MNIST



Small (28 X 28)

Clear and well-defined

Lots of examples



Fashion MNIST vs Inception V3

Fashion MNIST



Convolution

Pool


Concat

Dropout

Dense

Softmax

Inception



Inception Training Issues

Needs lots of data

Takes several weeks

1.2 million images

Big clusters of systems



Transfer Learning Steps

Use pre-trained model


- Leverage power of model
- Feature detection

Replace classifier with our classifier

Train classifier on our classes



Transfer Learning Retrain



Demo



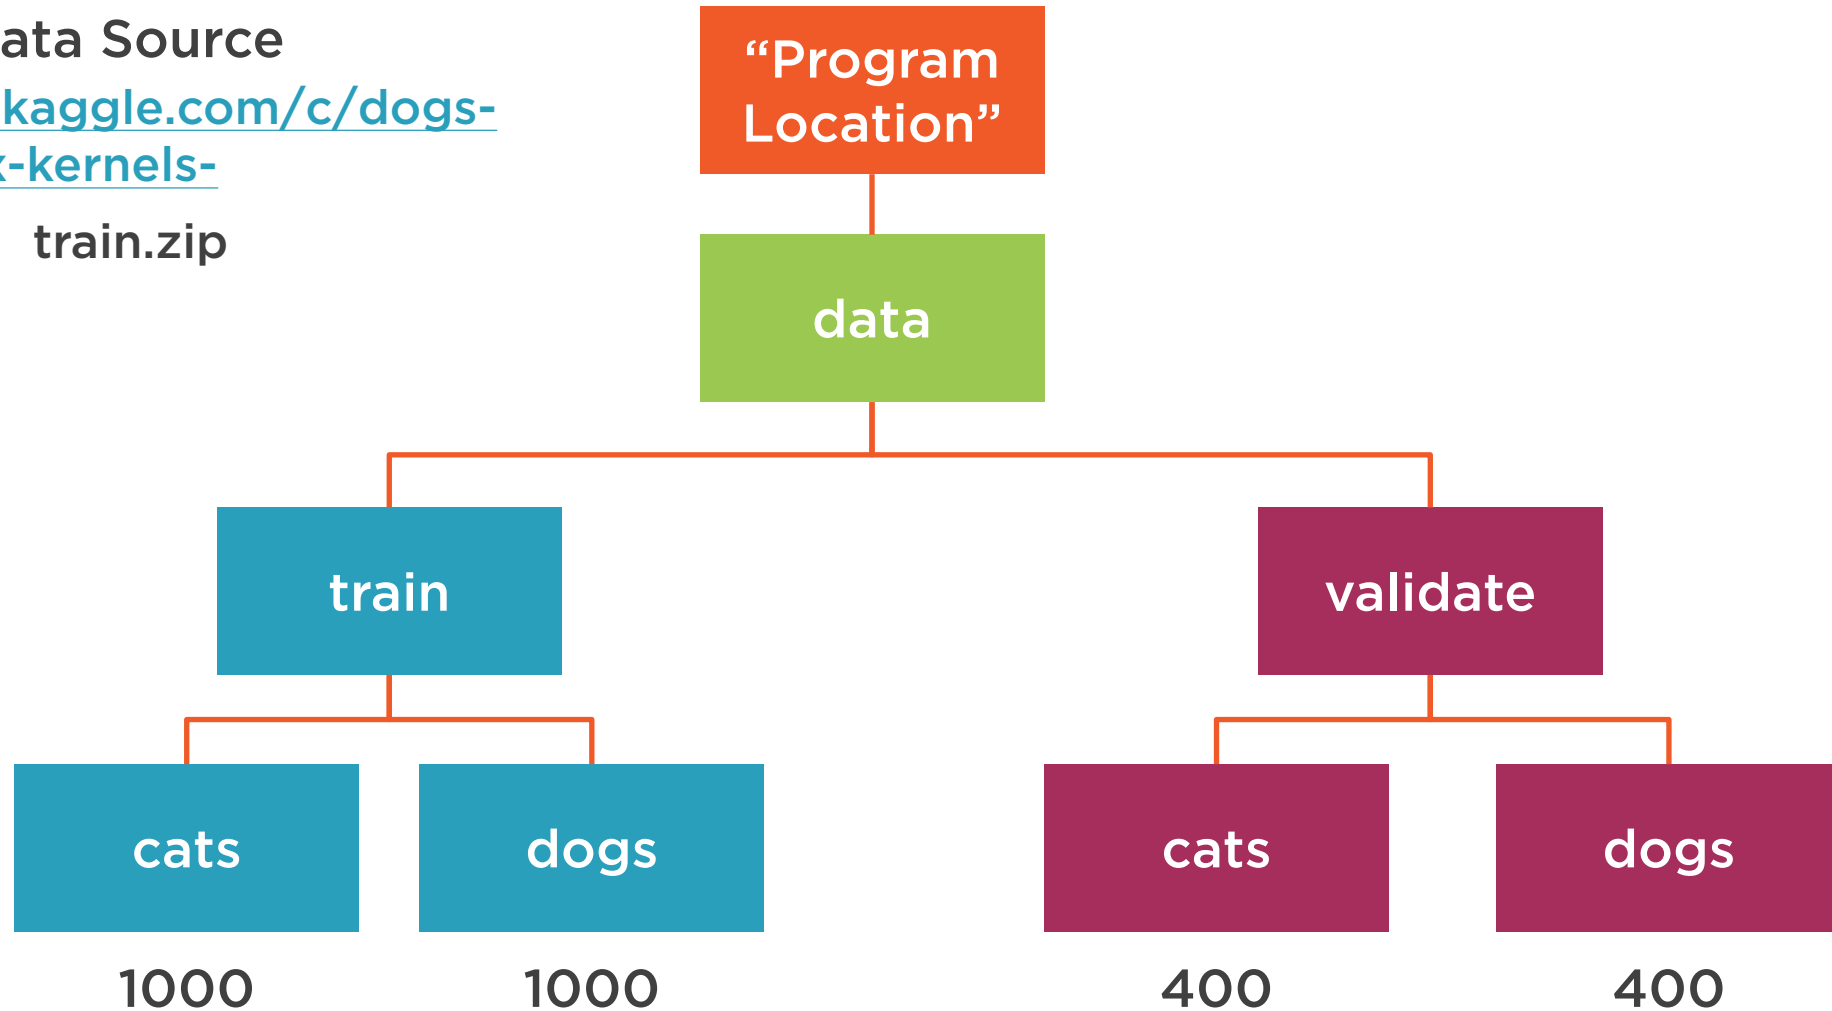
Transfer Learning



Data Folder Structure

Data Source

<https://www.kaggle.com/c/dogs-vs-cats-redux-kernels-edition/data> train.zip



More Than Static Images

Video

Signal Processing

Natural Language Processing

Combined with RNNs



Summary



Reviewed image data issues

- Lots of weights
- Translation Invariance

CNNs solved issues

CNNs build Feature Maps

Built FashionMNIST model from scratch

Utilized Transfer Learning

CNNs easy with Keras

